



**Afi**

Escuela  
de Finanzas

# Regresión Avanzada

## Máster en Data Science y Big Data en Finanzas

**Javier Nogales – PhD Matemáticas**  
**Catedrático, Estadística e IO, UC3M**

[fcojavier.nogales@uc3m.es](mailto:fcojavier.nogales@uc3m.es) 2022

# Organization

- Subject organized in 4 topics
- 12 hours in total: 3 sessions
- Practical course: 50% basic concepts + 50% computer labs (using R)
- Evaluation: 100% final exercise (40% GLM + 60% this topic)

# Objectives

- Relax some of the assumptions in classical linear regression (normality, loss functions, etc.)
- Deal with the curse of dimensionality in high-dimensional problems
- Handle the R language for advanced regression (including caret package)

# Outline

## 1. General Loss Functions

## 2. Regression Tools in High Dimension

- Model Selection
- Regularization Methods
- Dimension Reduction
- Feature Selection

## 3. Logistic Regression

# 1. General Loss Functions

# Linear regression: a brief review

- Remember the usual framework in **Machine Learning / Statistics**:

$$\text{Data} = \text{Model} + \text{Noise}$$

- When we focus on one variable predicted by others:

$$y = g(x_1, \dots, x_k) + \text{Noise}$$

- Most widely used tool (approximation) with  $p$  variables:

$$y = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p + \epsilon$$

- This is the **Multiple Linear Regression** tool
- Note the variables in the true model,  $g$ , may be different in the approximation

# Linear regression: a brief review

- For the model  $y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \epsilon_i$

- Assumptions

①  $E(\epsilon_i) = 0 \quad \forall i$

②  $\text{Var}(\epsilon_i) = \sigma^2 \quad \forall i$

③  $E(\epsilon_i \epsilon_j) = 0 \quad \forall i \neq j$

④  $\epsilon_i \sim \text{Normal} \quad \forall i$

①  $E(y_i) = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_k x_{ip} \quad \forall i$

②  $\text{Var}(y_i) = \sigma^2 \quad \forall i$

③  $\text{Cov}(y_i, y_j) = 0 \quad \forall i \neq j$

④  $y_i \sim \text{Normal} \quad \forall i$

- Note (3) and (4) imply the observations are independent
- We do not need (4) for estimation, but for inference

# What happens if some assumptions do not hold?

- In practice, the output may not be continuous
- The errors may not be normal
- The relation between the output and the predictors may not be linear

**Generalized Linear Models (GLM):** try to extend linear models by relaxing previous assumptions in an unified way

**Non-linear regression:** specify a parametric non-linear (in parameters) relation and estimate the parameters using an optimization solver

**Regression in high-dimension:** OLS has high variance in high dimension, advanced tools try to reduce the variance while increasing the bias

**Non-parametric regression / machine learning:** a non-linear function is estimated (learned) that cannot be parametrized



# Non-linear Regression

- Are these models linear?

$$y_i = \beta_0 + \beta_1 x_{i1}^2 + \beta_2 x_{i2}^2 + \beta_3 x_{i1} x_{i2} + \epsilon_i$$
$$\log(y_i) = \beta_0 + \beta_1 x_{i1} + \epsilon_i$$

- Are these models linear?

$$y_i = \beta_0 + \beta_1 x_{i1}^{\beta_2} + \epsilon_i$$
$$y_i = \frac{\beta_0}{1 + \frac{x_{i1}}{\beta_1^2}} + \beta_2 x_{i2} + \epsilon_i$$

# Parametric Non-linear Regression

- Analyze non-linear relationships in the form:

$$y = F(\beta; x) + \text{Noise}$$


- We assume the parametric function  $F$  is known, except for the parameters  $\beta$
- Estimation is done through **nonlinear least-squares**:

$$\text{minimize}_{\beta} \quad \frac{1}{2} \sum_i (y_i - F(\beta; x_i))^2$$

- Optimization algorithms are needed (gradients and Hessian matrix)
- Logistic regression can be viewed as a particular case, and also it is a particular case of GLM

# General Loss Functions

- Residuals in OLS can be written as  $\sum_i (y_i - x_i^T \beta)^2 = \|y - X\beta\|_2^2$
- Hence, other versions if we change the norm (loss function)
- For instance, loss function in OLS is  $\rho(x) = x^2$



# Least-absolute or $L_1$ regression

- The loss function is  $\rho(x) = |x|$
- It gives less weight to large errors (outliers):

$$\text{minimize}_{\beta} \quad \|y - X\beta\|_1 = \sum_i |y_i - x_i^T \beta|$$

No explicit solution

- This approach is **robust** against outliers: the absolute value gives less weight to points far from the center than OLS

# Chebyshev regression

- Minimax approach:

$$\text{minimize}_{\beta} \quad \|y - X\beta\|_{\infty} = \min \max |y_i - x_i^T \beta|$$

No explicit solution

- This approach is also robust but in a minimax sense: it minimizes the worst residual
- Very conservative approach

# Robust regression

- A general robust framework:

$$\text{minimize}_{\beta} \quad \sum_i \rho\left(\frac{y_i - F_i(\beta; \mathbf{x}_i)}{\sigma_i}\right),$$

where  $\sigma_i$  is a scale parameter (perhaps not included) and  $\rho$  is a loss (error) measure

# Robust regression: particular cases

- Least squares:  $\rho(x) = \frac{x^2}{2}$
- $L_p$  estimators:  $\rho(x) = |x|^p$ , with  $1 < p < 2$
- Lorentz estimator:  $\rho(x) = \log(1 + \frac{1}{2}x^2)$
- Truncated least-squares estimator:

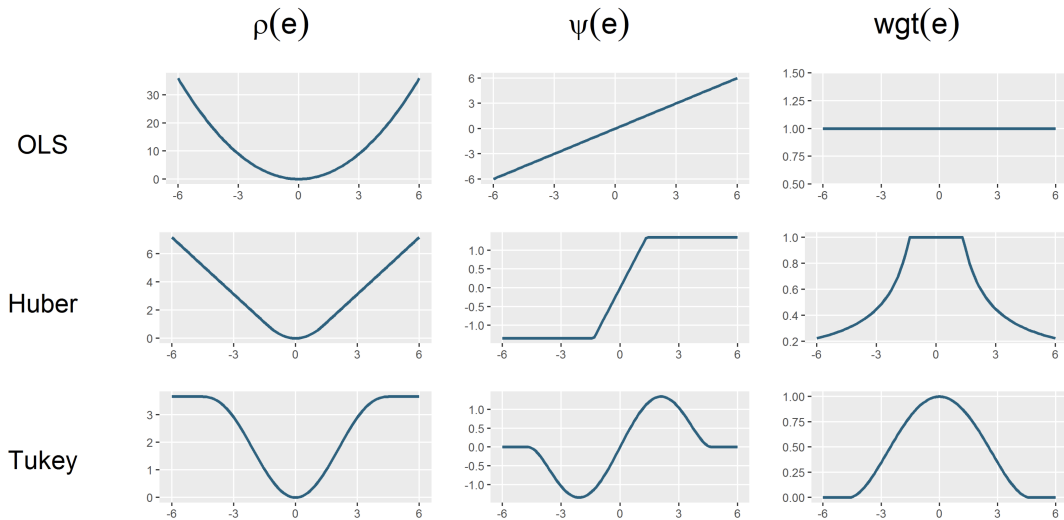
$$\rho(x, k) = \begin{cases} x^2, & |x| < \sqrt{k}, \\ k, & \text{otherwise} \end{cases}$$

- Huber function, etc.

$$\rho(x, k) = \begin{cases} \frac{1}{2}x^2, & |x| < k, \\ k|x| - \frac{1}{2}k^2, & \text{otherwise} \end{cases}$$

- Others

# Robust regression: Loss, Influence, and Weight Functions





## 2. Regression Tools in High Dimension

# Advanced Regression

How to improve simple linear models when the dimension is high?

- Try to improve the interpretability while attaining good predictive performance
- Need to replace least squares with some alternative fitting tools
- Need to balance prediction accuracy versus model interpretability (feature selection)

# Redundant information, collinearity and overfitting

- To understand the impact of high dimension on estimates and predictions, we need to understand first the differences between redundant information, collinearity and overfitting
- **Redundant information:** no relevant/valuable information is added to explain or predict better
- We tend to think more data is better, but this is usually wrong:

More data = more information + much more noise

- Redundant information implies collinearity and overfitting, but what's that?

# Collinearity

- Phenomenon indicating the predictors are linearly related (redundant information in a linear way)
- Hence, it affects basically linear models
- In high dimension ( $p/n$  is large), collinearity appears almost sure
- That implies a close-to-singularity matrix  $X'X$  (large condition number)
- That implies beta's are estimated with high noise
- That implies confusing and misleading results about the effects: non-reliable t-tests for parameters but reliable F-test, non-significant parameters in multiple regression but significant ones in simple regressions, etc.
- In linear models, collinearity implies overfitting, but what's that?

# Overfitting

- Condition indicating the model is focusing on noise rather than information (more parameters than needed)
- In practice, overfitting causes very good predictions in past data, but bad predictions in new data
- But how harmful is collinearity and overfitting in linear models for prediction?
  - Not harmful for prediction if new testing samples are within the range of training ones
  - Otherwise, there is extrapolation risk, and no limits in prediction error

# How harmful is overfitting in linear models?

- Estimation by OLS is not reliable (bad explanation):

$$E(\text{Var}(\hat{\beta})) = \sigma^2 \text{trace}((X^T X)^{-1})/p$$

- Prediction of output is reliable within the range of historical data (in-sample):

$$E(\text{Var}(\hat{y})) = E(\text{Var}(X\hat{\beta})) = \sigma^2 \text{trace}(X(X^T X)^{-1} X^T) = \sigma^2 p/n$$

- But maybe not reliable if predictors are outside that range:

$$\text{Var}(\hat{y}_0) = \text{Var}(x'_0 \hat{\beta}) = \sigma^2 (x'_0 (X^T X)^{-1} x_0)$$

- Hence, how can we estimate with some accuracy  $\beta$ ? To explain
- And, are all the  $p$  variables really needed to predict  $y$ ? To predict better

# To explain or to predict?

In regression/classification, there are three sources of uncertainty:

- The error in the coefficients when the linear approximation is true (**estimation error**)
- The error in the linear approximation when the true model is non-linear, or contains other variables (**model bias**)
- The noise in the DGP:  $\text{Data} = \text{Model} + \text{Error}$  (**irreducible error**)
- Error decomposition:  $\text{Prediction Error} = \text{Data} - \widehat{\text{Model}}$

$$(\text{Prediction Error})^2 = \sigma^2 + \text{Bias}^2 + \text{Var}$$

# To explain or to predict?

$$(\text{Prediction Error})^2 = \sigma^2 + \text{Bias}^2 + \text{Var}$$

- **Statistics:**


- Focus on minimizing Bias (by assuming knowledge about population, DGP)
- Hence, able to obtain formulas for Var that **provides explanation** (inference, effects of predictors on response)
- The Var can be large in practice

- **Machine Learning:**

- Focus on minimizing  $\text{Bias}^2 + \text{Var}$
- No assumptions needed (discover knowledge), hence no formulas and no explanation
- But **good prediction** performance



# Predictions: Bias vs Variance



## Prediction error of linear models

- General decomposition:  $(\text{Prediction Error})^2 = \sigma^2 + \text{Bias}^2 + \text{Var}$
- For a linear model  $y = x^T \beta + \epsilon$  with  $p$  variables (features) and  $n$  observations
- Assume no bias: we are using the same variables as in the population
- Prediction error in training set:

$$(\text{Prediction Error})^2 = E((y - \hat{y})^2 | x) = \sigma^2 + 0 + \sigma^2 \frac{p}{n}$$

It does not depend on collinearity, but increases with  $p$

- Prediction error in testing set:

$$(\text{Prediction Error})^2 = \sigma^2 + 0 + \sigma^2 (x_{\text{test}}^T (X_{\text{train}}^T X_{\text{train}})^{-1} x_{\text{test}})$$

That implies reliable predictions within the range of historical data ( $x_{\text{test}} \in X_{\text{train}}$ )

# Overfitting and underfitting in practice

- Consider the following true data generating process (DGP):

$$y = x_1\beta_1 + \cdots + x_p\beta_p + \epsilon = x'_{\text{true}}\beta + \epsilon$$

where  $E(\epsilon) = 0$  and  $\text{Var}(\epsilon) = \sigma^2$ . The corresponding OLS estimator is denoted by  $\hat{\beta}_{\text{true}}$

- Overfitting**: the model is estimated with more variables than needed ( $q > p$  variables):

$$y = x_1\beta_1 + \cdots + x_q\beta_q + \epsilon = x'_{\text{over}}\beta_{\text{over}} + \epsilon$$

If  $\hat{\beta}_{\text{over}}$  denotes the OLS estimator, then the prediction is unbiased but with larger variance:


$$E(x'_{\text{over}}\hat{\beta}_{\text{over}}) = x'_{\text{true}}\beta_{\text{true}}, \quad \text{Var}(x'_{\text{over}}\hat{\beta}_{\text{over}}) \geq \text{Var}(x'_{\text{true}}\hat{\beta}_{\text{true}})$$

- Underfitting**: the model is estimated with less variables than needed ( $q < p$  variables). Then, the prediction is biased but with smaller variance, i.e.

$$E(x'_{\text{under}}\hat{\beta}_{\text{under}}) \neq x'_{\text{true}}\beta_{\text{true}}, \quad \text{Var}(x'_{\text{under}}\hat{\beta}_{\text{under}}) \leq \text{Var}(x'_{\text{true}}\hat{\beta}_{\text{true}})$$

# Overfitting

When a model fits or predict very well the training data but bad the testing data  
( $p$  is large compared with  $n$ )



# Regression Tools in High Dimension

- Variable selection
- Regularization (shrinkage estimation)
- Dimension Reduction

- **Model Selection**

# Model Selection

Assume we have a high-dimensional regression problem:

$$y = \beta_0 + \beta_1 x_1 + \cdots + \beta_p x_p + \varepsilon = X\beta + \varepsilon$$

where  $p/n$  is large

Two main questions:

- How can we estimate with some accuracy  $\beta$ ?
- Are all the  $p$  variables needed to predict  $y$ ?

# Model Selection

- Strong objective: identify **the** subset of the  $p$  variables that is **really associated** to the response  
This may be prohibitive: if  $p = 50$ , then more than  $10^{15}$  models to try!
- Weak objective: identify **a** subset of the  $p$  variables that is **relevant** to the response  
But, what is the meaning of relevant?
- Some quality measures: adjusted  $R^2$ , Akaike information criterion (AIC), Bayesian information criterion (BIC), based on cross-validation, based on fitting in testing set, . . .



# Model Selection: AIC

- For a given estimate  $\hat{\beta}_d$  of the model  $y = X_d\beta_d + \varepsilon_d$ , where  $X_d$  represent a (column) subset of  $X$ , let  $RSS_d$  be the residual sum of squares:

$$RSS_d = \|y - X_d\hat{\beta}_d\|^2$$

and  $\hat{\sigma}_d^2 = \text{var}(\varepsilon_d)$  is the residual variance

- We need to penalize RSS to avoid overfitting
- Hence, instead of focusing on RSS, we focus on

$$AIC_d = \frac{1}{n\hat{\sigma}_d^2} (RSS_d + 2d\hat{\sigma}_d^2)$$

- Choose model with smallest  $AIC_d$

## Model Selection: BIC

- The BIC has a heavier penalty, hence select a smaller number of variables

$$\text{BIC}_d = \frac{1}{n}(\text{RSS}_d + \log(n) \, d \, \hat{\sigma}_d^2)$$

- Choose model with smallest  $\text{BIC}_d$

# Model Selection: Cross-validation

- Leave-one-out cross-validation:
  - split the sample in two parts, one with one observation (test) and the other with the rest (train)
  - fit the model with the training set, validate the model with the test set
  - repeat the process and compute the validation error using test observations
- $k$ -Fold Cross-Validation: similar, but randomly divide the sample in  $k$  groups, where the training set uses  $k - 1$  folds and the test set uses the remaining one

$k = 5$  or  $k = 10$  is typically used

- Fewer assumptions than AIC or BIC. Indeed, AIC is asymptotically similar to leave-one-out CV
- But more computational demanding

# Model Selection: testing set

- Similar to cross-validation:
  - **randomly** split the sample in two parts: one with (let's say) 80% observations (training set) and the other with the rest (testing set)
  - fit the model with the training set, validate the model with the testing set
  - repeat the process many times, computing each time fitting measures (like  $R^2$ ) in the testing set
- Computational demanding

# Model Selection: Subset Selection

## Best subset selection:

- Fit a regression for each possible combination of the  $p$  predictors:  $2^p$  combinations
- Fix a relevant criterion (AIC, BIC, cross-validation, ...)
- Choose the combination with smallest criterion

# Model Selection: Stepwise Selection

- Best subset selection is computationally infeasible for  $p > 40$
- Forward Stepwise Selection:
  - Start with a model with no predictors, with the corresponding quality measure (AIC, BIC, cross-validation, . . .)
  - Add one predictor each time, and select that with best fit. Compute the corresponding criterion for the best fit (AIC, BIC, cross-validation, . . .)
  - Add one (remaining) predictor each time and repeat the process
  - Select the model with smallest criterion
- This is an heuristic: only  $1 + p(p + 1)/2$  models considered (quadratic instead of exponential)

# Model Selection: Stepwise Selection

- **Backward Stepwise Selection:**
  - Start with the full model, with the corresponding quality measure (AIC, BIC, cross-validation, . . .)
  - Remove one predictor each time, and select that with best fit. Compute the corresponding criterion for the best fit (AIC, BIC, cross-validation, . . .)
  - Remove one (remaining) predictor each time and repeat the process
  - Select the model with smallest criterion
- This is again an heuristic: only  $1 + p(p + 1)/2$  models considered
- Not feasible when  $n < p$
- **Hybrid approaches:** variables are added and removed to the model sequentially (stepwise regression)
- Same ideas can be used to other regression models like logistic regression (where the deviance plays the role of RSS)

- **Regularization Methods**



# Regularization Methods

- How can we estimate with some accuracy  $\beta$ ?
- Main idea: penalize coefficient estimates, i.e. shrink them to 0
- With this framework, no need to select variables previously
- Main tools: Ridge, Lasso, Elastic Net, etc.

# Ridge Regression

- **Ridge regression**: used in high dimension to mitigate overfitting (even if  $p > n$ )
- Also known as Tikhonov regularization: ill-conditioned problems


$$\text{minimize} \quad ||y - X\beta||_2^2 + \rho||\beta||_2^2$$

where  $\rho$  is a tuning parameter, to be calibrated separately

- Explicit solution:  $\hat{\beta} = (X^T X + \rho I)^{-1} X^T y$
- It adds some bias to the estimation to reduce a lot the variance: better MSE than OLS
- It is better the data matrix  $X$  is centered previously (no estimation of  $\beta_0$ , we do not want to shrink it). Then,  $\hat{\beta}_0 = \bar{y}$
- It is also better to standardize the data, in order to make the estimation scale-invariant
- Low computational cost, good prediction accuracy, but dense solution (no variable selection)

# Ridge Regression

- Ridge regression coefficients as  $\log \rho$  increases



- Note ridge tends to OLS when  $\rho$  is small, and tends to 0 when  $\rho$  is large

# Ridge Regression: some theoretical results

- It is a biased estimator:

$$\text{Bias}^2(\rho) = \|E(\hat{\beta}_{RR}(\rho)) - \beta^*\|^2 = \|((X^T X + \rho I)^{-1}(X^T X) - I)\beta^*\|^2 \neq 0$$

- It has better variance than OLS:

$$\text{Var}(\rho) = \sigma^2 \text{trace} (X^T X + \rho I)^{-1}(X^T X)(X^T X + \rho I)^{-1} \leq \text{Var}(\hat{\beta}_{OLS}) = \sigma^2 \text{trace} (X^T X)^{-1}$$

- MSE decomposition:  $\text{MSE}(\hat{\beta}_{RR}(\rho)) = \text{Bias}^2(\rho) + \text{Var}(\rho)$

- Theorem (Theobald, 1974): there exists  $\rho > 0$  such that

$$\text{MSE}(\hat{\beta}_{RR}(\rho)) < \text{MSE}(\hat{\beta}_{OLS})$$

# Ridge Regression: a Bayesian view

- Up to now,  $\beta$  was fixed but unknown
- **Bayesian view:** treat it as a random quantity
- Before observing  $X$  and  $y$ , we assume a prior distribution for  $\beta$ ,  $p(\beta)$
- After observing  $X$  and  $y$ , we can compute the posterior distribution for  $\beta$  as

$$p(\beta|X, y) \propto p(\beta) p(X, y|\beta)$$

- If  $p(\beta) \sim \mathcal{N}_p(0, 1/(2\rho^2)I)$  (parameters are close to 0), then the mean of the posterior distribution is equivalent to ridge regression

# The Lasso

- **Lasso regression**: used in high dimension to mitigate overfitting (even if  $p > n$ )
- $L_1$  regularization: **sparse solutions**

$$\text{minimize}_{\beta} \quad \frac{1}{2} \|y - X\beta\|_2^2 + \rho \|\beta\|_1$$


- No explicit solution: non-differentiable problem
- It adds some bias to the estimation to reduce a lot the variance: better MSE than OLS
- Attains sparsity (model selection at the same time)
- Again, it is convenient the data matrix  $X$  is centered previously (no estimation of  $\beta_0$ , we do not want to shrink it). Then,  $\hat{\beta}_0 = \bar{y}$
- State-of-the-art tool in **Big Data Analytics**

# The Lasso

- Many ways to estimate the lasso regression: solving a quadratic optimization problem, solving the original (but non-differentiable) problem, using the LARS, using coordinate descent, . . .
- With non-prior information, ridge regression attains less variance than lasso (with similar bias)
- If real model is sparse, then lasso performs better
- Take care: if  $p > n$ , the Lasso selects at most  $n$  variables (which can be small)
- **Bayesian view:** if the prior for  $\beta$  is a Laplace distribution (some parameters are 0), then the posterior mean is equivalent to the Lasso

# The Lasso

- Lasso regression coefficients as  $\log \rho$  increases



- Note lasso tends to OLS when  $\rho$  is small, and tends to 0 when  $\rho$  is large, but in a sparse way



# Cardinality constraints or $L_0$ regression

- Based on 0-norm penalty:

$$\text{minimize}_{\beta} \quad \frac{1}{2} \|y - X\beta\|_2^2 + \rho \|\beta\|_0$$

where  $\|\beta\|_0$  is the number of non-zero elements in  $\beta$

- $\|\cdot\|_0$  is not really a norm...
- This formulation is equivalent to best subset selection
- Can improve computational efficiency by using good MIP techniques (optimization)
- Not quite stable in practice

# Regularization: Elastic Net


- Based on 1 and 2-norm penalties:

$$\text{minimize}_{\beta} \quad \frac{1}{2} \|y - X\beta\|_2^2 + \alpha \rho \|\beta\|_1 + \alpha(1 - \rho) \|\beta\|_2^2$$

- The 1-norm controls sparsity
- The 2-norm stabilizes the regularization path
- But we need to calibrate two parameters...
- Recommended packages for regularization tools: glmnet (in R), scikit.learn (in Python)

# Regularization: Tuning

- How to optimize the hyper-parameters  $\rho$ ?
- Perform cross-validation and select the parameter by minimizing sum of squared residuals



## Regression in high dimension: some conclusions

- Model selection tools are also dimension-reduction tools, but without changing the input. Usually they have smaller bias but larger variance
- Model selection focuses more on interpretation, less on prediction
- Regularization tools do not change the input, but change (bias) the estimator. Usually they have larger bias but smaller variance
- Lasso performs a bit worse than Ridge (in terms of prediction error) but provides model selection (better interpretation)

- **Dimension Reduction**

# Dimension Reduction Tools

- Principal Component Regression (**PCR**): reduce dimension in  $X$  using just a few PCs, and then perform the regression
- Partial Least Squares (**PLS**): similar in spirit to PCR but using (besides  $X$ ) the information in  $y$  to reduce the dimension  
Widely used in chemometrics, especially when  $p \gg n$
- PLS, PCR, and Ridge regressions perform similarly  
Ridge regression maybe preferred because simplicity

# Principal Component Regression

- **Principal component regression**: instead of fitting a high-dimensional regression model:

$$y = X\beta + \varepsilon, \quad \text{where } \beta \in \mathcal{R}^p$$

apply PCA:  $Z = XA_r$  with  $r \ll p$  and fit

$$y = Z\beta + \varepsilon, \quad \text{where now } \beta \in \mathcal{R}^r$$

- Because matrix  $Z$  is orthogonal, PCR is a sum of univariate regressions
- We need to standardize previously  $X$  and  $y$

# Partial Least Squares

- **Partial Least Squares**: similar in spirit to PCR
- PLS also looks for linear combinations of the predictors, but using  $y$  in addition to  $X$  to find the optimal combination
- In other words, PCR identifies directions in an unsupervised way whereas PLS uses the response to supervise the identification of the principal components
- I.e. PLS tries to find directions that help explain both the response and the predictors
- In general, PLS has better prediction performance than PCR, but not always
- We need to standardize previously  $X$  and  $y$



# Partial Least Squares: Procedure

- After standardizing the  $p$  predictors, the weights for the first linear combination are computed through the slopes of the simple regressions between  $y$  and  $X_j$  (instead of using the eigenvectors, as in PCA)
- That implies the weights of the linear combination are proportional to the correlation between  $y$  and  $X_j$
- Hence, PLS gives the highest weight to the most correlated variable with the response
- Subsequent directions are found by taking residuals and then repeating the above step

## Dimension-reduction tools: some conclusions

- PCR and PLS are dimension-reduction tools that change the input, hence they are less interpretable
- PCR and PLS perform roughly similar to Ridge regression, so the last one is usually preferable (much simpler and smoother)

# Optimization of hyper-parameters

Selecting the tuning parameter:

- How can we select the number of variables in step-wise regression, or the penalty parameters in regularization tools, or the number of directions in PCR or PLS?
- Cross-validation: choose a grid of those values, and compute the cross-validation error (quality measure) for each value of the grid
- Select the value with smallest cross-validation error

This is the usual way, but computationally expensive

- **Feature Selection**

# Feature Selection

- Remember the usual framework in **Machine Learning / Statistics**:

$$\text{Data} = \text{Model} + \text{Noise}$$

- Supervised Learning with real variables/features causing the target:

$$y = g(x_1, \dots, x_k) + \text{Noise}$$

- In Statistics or Machine Learning a set of  $p$  variables is used instead:  $x_1, \dots, x_p$
- Feature Selection**: out of these  $p$  features, how to select the most promising ones
- Ideally, we should select the  $k$  real ones, but impossible in practice

# Feature Selection

We can use previous knowledge:

- To perform  $p$  simple regressions and select the features with the highest correlation (or beta) with the output
- To perform one multiple regression over all the  $p$  variables and select the significant ones
- To consider the same multiple regression and focus on the highest t-values (variable importance)
- To apply stepwise regression
- To apply Lasso or Elastic Net

# Feature Selection

We can use new knowledge:

- Recursive Feature Elimination: based on backward selection and variable importance (using LM or RF, etc.)
- Random Forest: variable importance
- Others: foci, boruta, etc.

# 3. Logistic Regression

From Regression to Classification



# Logistic Regression

- Extension of classical multiple regression where the response was a continuous variable
- Now the response is a qualitative variable, usually binary, or the associated proportions
- Hence, logistic regression can be viewed as a particular case of
  - Non-linear regression
  - GLM (logit link)
  - Classification (linear classifier)
- The predictors can be either continuous or categorical
- Three types:
  - Binary: only two possible outcomes
  - Multinomial: more than two outcomes, no ordering
  - Ordinal: more than two outcomes, with a natural ordering

# Logistic Regression


- Process to predict a categorical/qualitative variable as a function of explanatory variables (predictors)
- Related with **Statistical Classification**: first predict the probability of each category, then predict the category
- Questions answered by Logistic Regression:
  - Which predictors are associated with the response?
  - What is the relationship between the response and the predictors?
  - Is a linear approximation valid?
  - How accurate is the prediction?
- Based on conditional distribution of the response given the predictors

# Applications

- Credit scoring, credit risk
- Stock trading
- Electronic fraud detection
- Spam filtering
- DNA sequence data to detect disease causing
- Probability of a heart disease in the next five years

# Binary Logistic Regression

- Imagine we have just two groups: binary regression.  
Then, the possible  $y$  labels are 0 or 1 (first and second group, respectively).
- Instead of modeling  $y$  as a continuous variable and then  $E(y|x_1, \dots, x_p) = \beta_0 + \beta^T x$




# Binary Logistic Regression

- Better to consider  $y$  is binary and then model

$$E(y|x_1, \dots, x_p) = p = P(y = 1|X = x) = F(\beta_0 + \beta^T x)$$

where  $F$  is a continuous function between 0 and 1



# Binary Logistic Regression

- The quantity  $\frac{p}{1-p}$  is called the **odds**
- An odds ratio of 1 indicates there is no association between the response and predictor
- A small value indicates very low probability that  $y = 1$ , whereas a large value indicates very high probability
- Usually, we impose that the log-odds or  $\text{logit}(p)$  is linear:

$$\text{logit}(p) = \log \frac{p}{1-p} = \beta_0 + \beta^T x$$

- In this case, a negative value indicates most likely  $y = 0$ , whereas a positive value indicates most likely  $y = 1$  (more symmetry)
- If we increase an x-variable by one unit, the logit is increased by the corresponding  $\beta$

# Binary Logistic Regression

- Hence,  $p = F(\beta_0 + \beta^T x) = \frac{\exp(\beta_0 + \beta^T x)}{1 + \exp(\beta_0 + \beta^T x)} = \frac{1}{1 + \exp(-\beta_0 - \beta^T x)}$
- Note  $P(y = 0|X = x) = 1 - p = \frac{1}{1 + \exp(\beta_0 + \beta^T x)}$
- Need optimization to estimate the  $\beta$ 's (based on the MLE of  $y$  given  $x$ )
- Non-linear LS could be used, but MLE has better statistical properties (GLM)

# Binary Logistic Regression

- MLE:  $\max \prod_{i=1}^n p_i(x_i|\beta)^{y_i} (1 - p_i(x_i|\beta))^{1-y_i}$
- The multivariate distribution of  $X$  is not so important: explanatory variables can be non-Gaussian, contain categorical variables, etc.
- What is important is that the logit be linear
- Taking logs:  $\max \sum_{i=1}^n (\beta_0 + \beta^T x_i) y_i - \sum_{i=1}^n \log(1 + \exp(\beta_0 + \beta^T x_i))$
- Use a convex optimization solver, problem is concave and differentiable
- After the estimation, for a given observation  $x$ , we can predict the associated probability:

$$\hat{p}(x) = P(y = 1) = \frac{\exp(\hat{\beta}_0 + \hat{\beta}^T x)}{1 + \exp(\hat{\beta}_0 + \hat{\beta}^T x)} = \frac{1}{1 + \exp(-\hat{\beta}_0 - \hat{\beta}^T x)}$$



# Binary Logistic Regression

- Note logistic regression is also a **linear classifier**: the decision boundary separating the two predicted classes is the solution of  $\hat{\beta}_0 + \hat{\beta}^T x = 0$
- That means, we are minimizing the mis-classification rate: we say  $y = 1$  when  $p \geq 0.5$  and  $y = 0$  otherwise. I.e. we say  $y = 1$  when  $\hat{\beta}_0 + \hat{\beta}^T x \geq 0$
- Moreover, the distance from a observation to the decision boundary is

$$\frac{\hat{\beta}_0 + \hat{\beta}^T x}{||\hat{\beta}||}$$

- Hence, class probabilities go towards the extremes (0 and 1) more rapidly when  $\hat{\beta}$  is larger
- If mis-classification errors are asymmetric, instead of classifying  $y = 1$  with rule  $\hat{p} > 0.5$ , change the 0.5 to be more conservative or risky

# Multinomial Logistic Regression

- General case:  $G$  groups. No order
- Examples: departments at a business (e.g., marketing, sales, HR), search engine used (e.g., Google, Yahoo!, MSN), and color (black, red, blue, orange), etc.
- Hence, labels  $y = \{0, 1, \dots, G - 1\}$
- First group is the control or reference group,  $y = 0$
- For the first group:

$$p_0 = P(y = 0 | X = x) = \frac{1}{1 + \exp(\beta_{0,1} + \beta_1^T x) + \dots + \exp(\beta_{0,G-1} + \beta_{G-1}^T x)}$$

# Multinomial Logistic Regression

- For the rest of the groups:

$$p_g = P(y = g|X = x) = \frac{\exp(\beta_{0,g} + \beta_g^T x)}{1 + \exp(\beta_{0,1} + \beta_1^T x) + \cdots + \exp(\beta_{0,G-1} + \beta_{G-1}^T x)}, \quad g = 1, \dots, G-1$$

- We need to impose  $p_0 + \dots + p_{G-1} = 1$
- Estimation is performed through optimization of the MLE
- Note in this case there are  $G - 1$  linear classifiers

# Multinomial Logistic Regression

- Take care: if we include explanatory variables that are not significant, then we may introduce a large bias in the estimation
- Hence, make a good variable selection before estimation, or estimate through shrinkage or regularization: add a penalty term of the form  $\rho ||\beta||$  in MLE
- The performance of the logistic regression model for  $G \geq 3$  is not so good. Better performance attained by other classification tools. But for  $G = 2$  the performance is good, as it allows for general predictors

# Ordinal Logistic Regression

- Again there are  $G$  groups, but now the order matters
- Examples: effectiveness on a scale of 1-5, levels of flavors for hot wings, medical condition (e.g., good, stable, serious, critical), etc.
- Now we have

$$\sum_{g=1}^{G^*} p_g = \frac{\exp(\beta_{0,g} + \beta^T \mathbf{x})}{1 + \exp(\beta_{0,g} + \beta^T \mathbf{x})},$$


where  $G^* \leq G$ , and  $p_1 \leq p_2 \leq \dots \leq p_G$

- Note this model is a cumulative sum of probabilities which involves just changing  $\beta_0$
- Estimation is performed through optimization of the MLE

# Logistic Regression: Classical Example

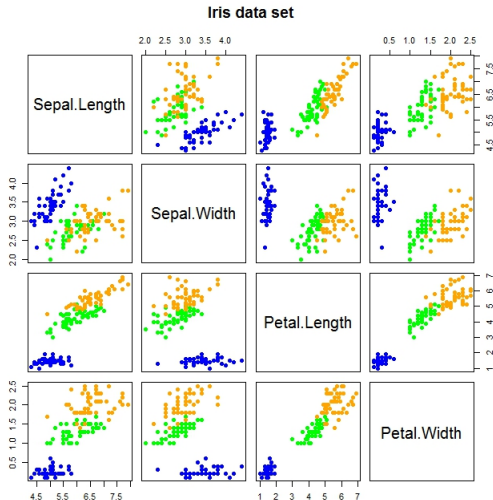
- Fisher's Iris data set: best known data set in classification
- Data collected by Edgar Anderson in the 30's to quantify the morphologic variation of Iris flowers of three related species
- It consists of 150 Iris flowers of three different species: setosa, versicolor, virginica
- Four types of measurements for each flower, the length and width of sepals and petals in centimeters, respectively
- Popularized by Ronald Fisher in 1936
- Based on the combination of the four features, Fisher developed a linear discriminant model to distinguish the species from each other

# Iris Flower



# Iris Scatter Plot

In blue, setosa Iris; in green, versicolor Iris; in orange, virginica Iris





## Iris example

- First linear classifier (versicolor respect to setosa):

$$8.1 + (-7 \ -7.2 \ 14 \ 18)x = 0$$

- Second linear classifier (virginica respect to setosa):

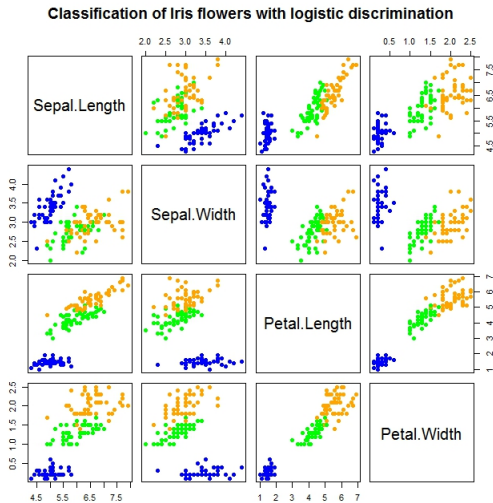
$$-34.5 + (-9.5 \ -13.9 \ 23.4 \ 36.3)x = 0$$

- Posterior probabilities for 120-th flower:

$$p_{\text{setosa}} = 0, p_{\text{versicolor}} = 0.08, p_{\text{virginica}} = 0.92$$

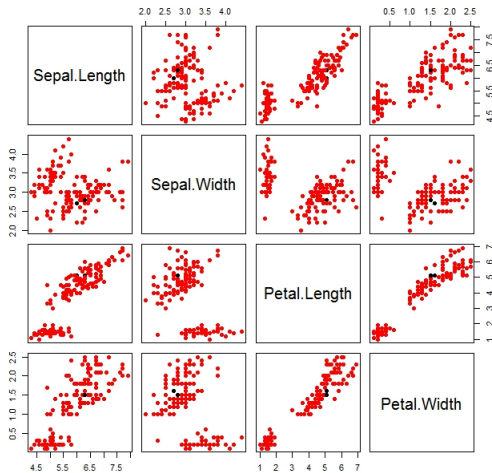
- Hence, predict observation 120-th is virginica
- But in general, how to select the threshold probability to classify?

# Iris: Logistic regression



# Iris: Logistic regression

Bad (in black) classifications for Iris flowers with logistic discrimination



# Classification: performance measures

- For a given threshold probability, we have the following table (confusion matrix):

	Classify in $P_1$	$\dots$	Classify in $P_G$
Belongs to $P_1$	$n_{11}$	$\dots$	$n_{1G}$
$\vdots$	$\vdots$	$\ddots$	$\vdots$
Belongs to $P_G$	$n_{G1}$	$\dots$	$n_{GG}$

where  $n_{ij}$  is the amount of observations coming from partition  $i$  that are classified in partition  $j$

- The **error** is then:  $\frac{n_{12} + \dots + n_{G,G-1}}{n}$ , the sum of off-diagonal elements
- The **accuracy** is the sum of the diagonal terms
- The confusion matrix should be computed using cross-validation or testing sets (not training ones)

# Performance measures

- The in-sample or training error may be **too optimistic**, due to overfitting
- We need more realistic errors:
- Use **cross-validation**: exclude from the sample one observation, estimate parameters to build the classification rule, classify the excluded observation, and finally check the error. Hence,  $n$  classifications are performed
- Use a **training** and a **validation set**: let's say, 80% of the data to estimate (train or learn) the model and the other 20% to validate. Repeat many times
- These are known as **out-of-sample** or **testing errors**
- Finally, note from confusion matrix other performance measures can be obtained: Accuracy, Kappa, Sensitivity (TPs), Specificity (TNs), etc.

## Iris example: logistic regression performance

- In sample performance:


predictions	setosa	versicolor	virginica
setosa	50	0	0
versicolor	0	49	1
virginica	0	1	49

- In-sample error = 1.3% (small, but optimistic)
- Cross-validation: error = 2% (a bit worse, but more realistic)
- Out-of-sample or testing: error = 3% (a bit worse, but more realistic)

# ROC curves

- Used for binary problems by varying the cutoff to assign groups (Bayes' rule)
- Compare false positive rates,  $P(\hat{y} = 1|y = 0)$ , with true positive ones,  $P(\hat{y} = 1|y = 1)$
- The true positive rate is also called the sensitivity and 1 minus the false positive rate is also called the specificity
- The receiver operating characteristic (ROC) curve is created by varying the cutoff from 0 to 1

# ROC curves





# AUC

- In previous graph, models 2 and 3 have higher true positive rates than model 1
- But difficult to compare models 2 and 3
- A useful summary of the overall quality is the area under the curve (AUC)
- An AUC of 0.5 indicates a random guessing while an AUC of 1 indicates perfect classification

## The end... Attained objectives

- Relax some of the assumptions in classical linear regression (normality, loss functions, etc.)
- Deal with the curse of dimensionality in high-dimensional problems
- Handle the R language for advanced regression (including caret package)

*“Far better an approximate answer to the right question, which is often vague, than an exact answer to the wrong question, which can always be made precise.”*

John W. Tukey



**Afi**

Escuela  
de Finanzas

---

© 2015 Afi Escuela de Finanzas. Todos los derechos reservados.