

# Neural Network Basics

José R. Dorronsoro  
Dpto. de Ingeniería Informática  
Escuela Politécnica Superior  
Universidad Autónoma de Madrid

## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

# What Is Machine Learning (ML)?

Machine Learning Basics   Multilayer Perceptrons

- Lofty definition: make machines learn!!!
  - Have to make “machines” and “learn” more precise
- The machines of ML: mathematical input–output processes that lend themselves to some form of (numerical) parameterization
- The learning process: adjust the machine’s parameters until a goal is reached
- New thing: “goal”?
  - At first sight, get something done
  - Ultimately, to minimize some error measure
- Summing things up: a ML process tries to find a concrete mathematical/algorithmic **input–output parameterized transformation** that **minimizes an error measure** by iteratively **adjusting the transformation’s parameters**




- In the middle of a possibly long process chain
- Before ML starts we must
  - Go from **raw to organized** data: accesing, gathering, cleaning, formatting, ...
  - Go from **organized to** (potentially) **informative** data: extracting basic and derived features
- After ML finishes and start applying a model, we must perform
  - Outcome **evaluation**: how good/actionable is
  - Outcome **exploitation**: collect, organize, act
  - **Individual model maintenance**: monitor performance, tune hyper-parameters
  - **Modeling life cycle maintenance**: discard old models, introduce new ones and **communicate** our work/results

- ML model types: **supervised, unsupervised**
- Supervised models:
  - Targets  $y^p$  are known and the model tries to predict or estimate them
  - These known targets guide, or **supervise**, model building
  - Main emphasis here
- Unsupervised models:
  - There are no predetermined or supervising outputs
  - But nevertheless the model is supposed to learn relations or find structure in the data
  - Often as a first step towards a supervised model
- Semisupervised models lie in between: some patterns have targets, some don't

- Problems (usually) to be solved by ML models: regression, classification
- In both patterns come in pairs  $(x, y)$ 
  - $x$ : inputs, predictors, features, independent variables
  - $y$ : target, response, dependent variable; numerical in regression, class labels in classification
- **Regression**: the desired output  $y$  is **regressed** into the inputs  $x$  to derive a model  $\hat{y} = f(x)$ 
  - We want  $y \simeq \hat{y}$  so having  $y - \hat{y}$  “small” is the natural goal
- **Classification**: inputs are derived from several classes  $C_1, \dots, C_K$ , to which labels  $\ell_k$  are assigned
  - The model now assigns a label  $\ell(x)$  to an input  $x$
  - If  $x$  is derived from  $C_k$  we want to have  $\ell(x) = \ell_k$
  - Here having  $\ell(x) - \ell_k$  “small” may not make sense

- This is a first, widely used “toy” problem
- We want to estimate the median of house values over an area from some information about it which we believe relevant
- Features  $x$ : several real estate–related variables of Boston areas
  - CRIM: per capita crime rate by town
  - RM: average number of rooms per dwelling
  - NOX: nitric oxides concentration (parts per 10 million)
  - AGE: proportion of owner-occupied units built prior to 1940
  - LSTAT: % lower status of the population
  - ...
- Target  $y$ : MEDV, median value of owner-occupied homes in \$1,000's

- This is a well known benchmark, moderately large classification problem
- We have  $28 \times 28$  grey scale image rasters with handwritten digits 0 to 9
- We want to build a classifier that assigns each raster to one of the 0 to 9 digits
- Features  $x$ :  $28 \times 28$  image matrices
- Target  $y$ : labels  $[0, 1, \dots, 9]$
- Examples (0-padded to  $32 \times 32$ ):





- Often individual models are selected through (ideally optimal) **parameter sets**
  - The parameters (weights)  $w \in R^M$  select a concrete  $f$  in a model family  $\mathcal{F}$
- **Parametric** models have a fixed functional form  $f(x) = f(x; W)$ 
  - Simplest example: **linear regression**, where  $x$  has dimension  $d$ ,  $M = d + 1$  and  $W = (w_0, w)$

$$f(x; w_0, w) = w_0 + \sum_{j=1}^d w_j x_j = w_0 + w \cdot x$$

- **Semi-parametric** models: also use weights but without a predefined functional form
  - MLPs but also RF or GBR
- **Non parametric** models do not use weights nor follow any broad functional form
  - Nearest Neighbor models

- There are some initial questions when working with models from a given family  $\mathcal{F}$ :
  - Which assumptions do they make?
  - How do they work on the input features?
  - How we do build them?
- In turn, these questions lead to another two:
  - How do we select the best model from the given family for the problem at hand?
  - How do we control the model building procedure?
- All of them address fundamental issues that require a **moderately deep** understanding of what is going on under the model's hood
- This understanding is usually framed in **mathematical language**

## 1 Machine Learning Basics

- Basic Modeling
- **Regression Basics**
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

## Machine Learning Basics Multilayer Perceptrons

- In general we have a sample  $S = \{x^p, y^p\}$ ,  $1 \leq p \leq N$ , with  $x^p$  the **features** and  $y^p$  the **targets**
- We want to build a model  $\hat{y} = f(x)$  so that  $\hat{y}^p = f(x^p) \simeq y^p$ ;
  - I.e., we want to **regress**  $y$  to the  $f(x)$
- The concrete  $f$  is chosen within a certain family  $\mathcal{F}$ 
  - Examples here: linear regression, multilayer perceptrons (MLPs)
  - And also: SVMs, Random Forests (RF), Gradient Boosting (GB), nearest neighbor (NN)
- Natural option to ensure  $f(x^p) \simeq y^p$ : choose  $f$  to minimize the sample **Mean Square Error (MSE)**

$$\hat{e}(f) = \hat{e}_S(f) = \frac{1}{2N} \sum_{p=1}^N (y^p - f(x^p))^2$$

- This corresponds to the **population** error

$$\begin{aligned} e(w_0, w) &= \frac{1}{2} E_{x,y} [(y - f(x))^2] = \frac{1}{2} E[(y - f(x))^2 p(x, y) dx dy] \\ &= \frac{1}{2} \int (y - f(x))^2 p(x, y) dx dy \end{aligned}$$

- Thus, we may select a model by solving  $\hat{f} = \hat{f}_S = \arg \min_{f \in \mathcal{F}} \hat{e}_S(f)$ 
  - But this may be quite complicated mathematically and we must simplify it
- For a parametric or semiparametric  $f(x; W)$  we can write  $\hat{e}_S(f) = \hat{e}_S(f(\cdot; W)) = \hat{e}_S(W)$
- The problem to solve becomes

$$\hat{W}^* = \hat{W}_S^* = \arg \min_W \hat{e}_S(W), \text{ i.e., } \hat{e}_S(\hat{W}^*) \leq \hat{e}_S(W) \forall W$$

- For linear regression (LR) the sample error is

$$\hat{e}(w_0, w) = \frac{1}{2N} \sum_p (y^p - w_0 - w \cdot x^p)^2$$

- Minimizing the sample error  $\hat{e}_S(W)$  becomes a problem in mathematical **optimization**

- **Key assumption:**  $x$  and  $y$  are related as  $y = \phi(x) + n$  where
  - $\phi(x)$  is the **true** underlying function
  - $n$  is **additive noise** with 0 mean and finite variance  $\sigma_N^2$
- Our sample is just a particular instance of a deeper **sample generation process**
- Thus  $x, n$  are produced by **random variables**  $X, N$ 
  - And so is  $y$ , given by  $Y = \phi(X) + N$
- Moreover,  $X$  and  $N$  are **independent distributions** with densities  $q(x), \nu(n)$
- Thus,  $X$  and  $Y$  (or  $X$  and  $N$ ) have a joint density

$$p(x, y) = p(x, \phi(x) + n) = q(x) \nu(n) = q(x) \nu(y - \phi(x))$$

- We can decompose the MSE error of any model  $f$  as

$$\begin{aligned}
 2\text{mse}(f) &= E_{x,y}[(y - f(x))^2] = \int (n + \phi(x) - f(x))^2 q(x) \nu(n) dx dn \\
 &= \int (n^2 + 2n(\phi(x) - f(x)) + (\phi(x) - f(x))^2) q(x) \nu(n) dx dn \\
 &= \int n^2 \nu(n) dn + \int (\phi(x) - f(x))^2 q(x) dx + \\
 &\quad 2 \left( \int n \nu(n) dn \right) \left( \int (\phi(x) - f(x)) q(x) dx \right) \\
 &= \sigma_N^2 + E_x[(\phi(x) - f(x))^2]
 \end{aligned}$$

- Thus **for any model we have  $\text{mse}(f) \geq \sigma_N^2$  always**
  - Beware of zero regression errors!!
- And we should focus on achieving on  $f \simeq \phi$  (which we don't know!!)


- It is easy to see that the best  $f$  is simply  $f(x) = E_y[y|x]$ , for

$$E_y[y|x] = E_n[\phi(x) + n] = \int (\phi(x) + n) \nu(n) dn = \phi(x)$$

- Have we finished? In theory yes; in practice, not at all!!!
  - We do not know  $\nu$  and, thus, cannot compute the required integral
  - If we would have several  $M$  values  $y^j$  for any  $x$ , we could try
 
$$\hat{\phi}(x) = \frac{1}{M} \sum_1^M y^j$$
  - But this doesn't happen either
- Now we have two options:
  - Try to stretch the  $E[y|x]$  approach, for instance, by  **$k$ -NN Regression**
  - Forget about it and get back to get models  $f$  for which we hope  $f \simeq \phi$



- The left hand side shows the ideal situation but the right one is what we should expect in practice



- We can try to use the  $y^p$  values of several  $x^p$  close to a new  $x$
- We fix  $k$  neighbors  $x^{p_1}, \dots, x^{p_k}$  of  $x$  and estimate  $\hat{y} = \hat{y}(x)$  as

$$\hat{y}(x) = \frac{1}{k} \sum_{j=1}^k y^{p_j}$$

- $\hat{y}(x) = \hat{Y}_k^{NN}(x)$  is the  $k$ -**Nearest Neighbor (NN)** regressor
  - This can be refined to weighted versions, such as

$$\hat{y}(x) = \frac{1}{C_k(x)} \sum_{j=1}^k \frac{1}{\|x^{p_j} - x\|^2} y^{p_j}$$

with  $C_k(x) = \sum_{j=1}^k \frac{1}{\|x^{p_j} - x\|^2}$  a normalizing constant

- But we need that close predictors give predictions that are also close, and that **there are enough of them close by**
  - This is **very unlikely**

- Even for low dimensions and large samples, **the sample space is essentially empty**
- Assume we have 1,000  $d$ -dimensional  $x$  patterns whose features have values between 1 and 10
  - In dimension  $d = 1$  there are 100 patterns per unit length
  - But when  $d = 3$  we have just 1 pattern per volume unit
  - And if  $d = 6$  we have just 1 pattern per 1,000 units of volume
  - And in dimension 10 (not a big one nowadays) we have just ... !!!
- Thus, for most problems, **there never will be enough close points**
- As a consequence, to get  $k$  observations we may go too far away from  $x$  and the average will not be meaningful
- Therefore, unless we deal with violently non-linear problems, a simple linear model may be better than  $k$ -NN regression for moderate dimensions

- Assuming  $x \in R^d$ , recall the basic linear model is

$$f(x) = w_0 + \sum_1^d w_i x_i = w_0 + w \cdot x$$

- $w_0$  complicates notation; to drop it we center  $x$  and  $y$  so that  $E[x_i] = E[y] = 0$ ; then  $w_0 = 0$
- Then we are left with the simpler **homogeneous** model  
 $f(x) = w \cdot x$
- In practice we will always **normalize**  $x$ , for instance to have 0 mean and 1 standard deviation (std) on each feature
  - But not  $y$  if we may help it (or it is easy to reverse this)
- But: how do we find  $w$ ?

- Assume that features  $X$  and target  $Y$  are **centered**, i.e., have 0 means
- For 1-dimensional patterns  $x$  the LR model then becomes  $f(x) = w x$
- And the error is then the function

$$\begin{aligned}\hat{e}(w) &= \frac{1}{2N} \sum_{p=1}^N (w x^p - y^p)^2 = \frac{1}{2N} \sum_p (w^2 (x^p)^2 - 2x^p y^p w + (y^p)^2) \\ &= \frac{1}{2} \left( \frac{1}{N} \sum_p (x^p)^2 \right) w^2 - \left( \frac{1}{N} \sum_p x^p y^p \right) w + \frac{1}{2N} \sum_p (y^p)^2\end{aligned}$$

- Thus,  $\hat{e}(w) = aw^2 + bw + c$  with  $a > 0$  and it has obviously a minimum  $w^*$
- To find it we just solve  $\hat{e}'(w) = 0$

- We first compute  $\hat{e}'(w)$ , for which we have

$$\hat{e}'(w) = w \left( \frac{1}{N} \sum_p (x^p)^2 \right) - \frac{1}{N} \sum_p x^p y^p$$

- The optimal  $w^*$  solves  $\hat{e}'(w) = 0$  and is given by

$$w^* = \frac{\frac{1}{N} \sum_p x^p y^p}{\frac{1}{N} \sum_p (x^p)^2} = \frac{\frac{1}{N} X^t Y}{\frac{1}{N} X^t X} = \frac{\text{covar}(x, y)}{\text{var}(x)}$$

where  $X$  and  $Y$  denote the  $N \times d$  **data matrix** (vector here) and the  $N \times 1$  **target vector**

- Assume again that  $X$  and  $Y$  are centered
- The LR model becomes now  $f(x) = \sum_1^d w_i x_i = w \cdot x$
- If  $Y$  is the  $N \times 1$  **target** vector and we organize the sample  $S$  in a  $N \times d$  **data matrix**  $X$ , the sample mse is given by

$$\begin{aligned}\hat{e}(w) &= \frac{1}{2N} \sum_p (w \cdot x^p - y^p)^2 = \frac{1}{2N} (Xw - Y)^t (Xw - Y) \\ &= \frac{1}{2N} (w^t X^t X w - 2w^t X^t Y + Y^t Y)\end{aligned}$$

- We have now a **quadratic form**  $w^t A w + w^t b + c$ , with  $A$  a  $d \times d$  matrix,  $b$  a  $d$  vector and  $c$  a scalar
- Now we have to solve  $\nabla \hat{e}(w) = 0$ , i.e.,  $\frac{\partial \hat{e}}{\partial w_i}(w) = 0, 1 \leq i \leq d$
- It is easy to see that

$$\nabla \hat{e}(w) = \frac{1}{N} X^t X w - \frac{1}{N} X^t Y = \hat{R} w - \hat{b}$$

- The optimal  $\hat{w}^*$  must verify  $\nabla \hat{e}(\hat{w}) = \hat{R} \hat{w} - \hat{b} = 0$ , where

$$\hat{R} = \frac{1}{N} X^t X, \quad \hat{b} = \frac{1}{N} X^t Y$$

- Over the original, non-centered data matrix we have

$$\hat{R} = \frac{1}{N} (X - \bar{X})^t (X - \bar{X});$$

i.e.,  $\hat{R}$  is the **sample covariance matrix**

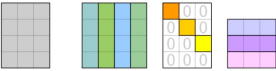
- If  $\hat{R}$  is **invertible**, we just solve the linear system  $\hat{R} \hat{w} = \hat{b}$  to get

$$\hat{w}^* = \hat{R}^{-1} \hat{b} = (X^t X)^{-1} X^t Y = \text{covar}(X)^{-1} \text{covar}(X, Y)$$

- We won't but we should write  $\hat{w}_S^*$ , as it is **sample-dependent**
- $\hat{R}$  is oftent inverted through the **Singular Value Decomposition (SVD)** of the data matrix  $X$



- The **Singular Value Decomposition** (SVD) of a (complex)  $m \times n$  matrix  $M$ ,  $m \geq n$ , is  $M = UDV^*$  where
  - $U$  and  $V$  are orthonormal (unitary) matrices with dimensions  $m \times n$ ,  $n \times n$  respectively
  - $D$  is an  $m \times n$  matrix, diagonal on its first  $n$  rows and with zeros in its last  $m - n$  ones
- Graphically, the SVD decomposition of  $M$  is



$$\begin{matrix} \mathbf{M} & = & \mathbf{U} & \mathbf{\Sigma} & \mathbf{V}^* \\ m \times n & & m \times m & m \times n & n \times n \end{matrix}$$

Taken from *Wikipedia*

- The columns of  $U$  and  $V$  are called the **left and right singular vectors**; the diagonal elements in  $D$  are the **singular values**
- We may interpret this image as  $M$  performing first a rotation with  $V^*$ , then feature dilations with  $\Sigma$  (our  $D$ ) and, finally, another rotation with  $U$ 
  - Actually both are a roto-reflections

- It is easy to see that we can also write  $M = \tilde{U}\tilde{D}V^*$ , with  $\tilde{U}$  having the first  $n$  columns of  $U$  and  $\tilde{D}$  the top diagonal block of  $D$ 
  - Assuming  $M$  to be real and with a slight abuse of language, we will write  $M = UDV^t$
- In our case we can write  $X = UDV^t$ , with dimensions  $N \times d$ ,  $d \times d$  and  $d \times d$  respectively
- Then  $X^tX = VDU^tUDV^t = VD^2V^t$  and  $(X^tX)^{-1} = VD^{-2}V^t$ 
  - $D^{-2}$  is a diagonal matrix with  $(D^{-2})_{ii} = 1/D_{ii}^2 = 1/\delta_i^2$
- The columns of  $U$  and  $V$  are unique up to a sign change
  - So different packages may give seemingly different  $U$  and  $V$

- Computing the covariance matrix has a  $O(N \times d^2)$  cost and invert it has a  $O(d^3)$  cost
  - For big data problems it may not be possible to solve analytically the normal equation  $\nabla \hat{e}(w) = 0$
- The simplest numerical alternative is **gradient descent**:
  - Starting from some random  $w^0$ , we iteratively compute

$$w^{k+1} = w^k - \rho_k \nabla \hat{e}_B(w^k) = w^k - \frac{\rho}{n_B} (X_B^t X_B w^k - X_B^t Y)$$

over a **mini-batch**  $X_B$  with  $n_B$  samples and batch error  $\hat{e}_B$

- Component wise:  $w_i^{k+1} = w_i^k - \rho_k \frac{\partial \hat{e}_B}{\partial w_i}(w^k)$
- $\rho_k$  is the **learning rate**
- If  $w^k \rightarrow w^*$ , then  $\nabla \hat{e}(w^*) = 0$ 
  - Since our problems have obviously minima, this should be enough

- We have just built our first neural network: the linear regression architecture can be seen as a multilayer perceptron (MLP) with no hidden layers
- Thinking of this as an MLP seems far-fetched, as we are missing the hidden layers and the much more complicated gradient computation
- But the layer architecture is there and we also have learning as iterative error minimization, i.e., **training**
- In fact, we will see that MLPs perform **linear regression on the last hidden layer outputs**
- And the process from the input to the last hidden layer produces an **enhanced feature representation** upon which the regression error should be better

- First option: **Root Square Error**  $RSE = \sqrt{\frac{1}{N} \sum (y^p - \hat{y}^p)^2}$
- OK, but how to compare this? We must always have a **base model** to benchmark our results
- Simplest model: a constant  $\hat{y} = w_0$ , which reduces to predict the mean  $\bar{y} = \frac{1}{N} \sum_1^N y^p$ , with square error

$$\frac{1}{N} \sum (y^p - \bar{y})^2 = \text{Var}(y)$$

- We can compare our model against this base model by computing

$$\frac{\sum (y^p - \hat{y}^p)^2}{\sum (y^p - \bar{y})^2} = \frac{\frac{1}{N} \sum (y^p - \hat{y}^p)^2}{\frac{1}{N} \sum (y^p - \bar{y})^2} = \frac{RSE^2}{\text{Var}(y)}$$

- The widely used  $R^2$  coefficient is simply  $R^2 = 1 - \frac{RSE^2}{\text{Var}(y)}$

- Our regression solution  $\hat{w}^* = (X^t X)^{-1} X^t Y$  won't work if  $X^t X$  is not invertible
  - For instance, when some features are highly correlated
  - We can at most say that  $X^t X$  is **positive semidefinite**
- We could fix this working instead with  $X^t X + \alpha I$  for some  $\alpha > 0$ 
  - This is then always **positive definite** and we can thus invert it
- To make this practical, note that  $\hat{w}^* = (X^t X + \alpha I)^{-1} X^t Y$  minimizes

$$e_R(w) = \frac{1}{2N} \sum_p (y^p - w \cdot x_p^p)^2 + \frac{\alpha}{2} \|w\|^2,$$

- This is the **Ridge Regression** problem
  - Our first example of **regularization**, a key technique in Machine Learning
  - **All ML models must be regularized in some way**
- Important issue: how to find the right choice for  $\alpha$ ?

- 1 We introduced **supervised** models
- 2 We have reviewed the essentials of the **linear regression model** (always the first thing to try)
- 3 We have considered model estimation as a problem on **error minimization**
- 4 We have seen how to build linear models **analytically and numerically**
- 5 We have seen how to **measure model fit**
- 6 We have introduced **regularization**

## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- **Bias, Variance and Cross Validation**
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs



- Remember: all the models  $f$  we will build are **sample dependent**, i.e.,  $f = f_S$
- With several **independent** samples  $S_1, \dots, S_M$ , it is natural to use as our best final model the average of their associated  $\hat{f}_{S_m}(x)$  models, i.e.,

$$\frac{1}{M} \sum_1^M \hat{f}_{S_m}(x) \simeq E_S[\hat{f}_S(x)] = \hat{f}_N(x)$$

- Here we take the (theoretical) expectation  $E_S[\hat{f}_S(x)]$  over all samples  $S$  of size  $N$
- $\hat{f}_N(x) = E_S[\hat{f}_S(x)]$  is our ideal **best model** and  $\phi(x) - \hat{f}_N(x)$  is its **bias**
- The **variance** of the  $\hat{f}_S(x)$  estimates is then

$$V_N(x) = E_S [(\hat{f}_S(x) - \hat{f}_N(x))^2]$$


- Recall that our goal is to estimate  $\phi$  in our regression model  $y = \phi(x) + n$
- Ideally we would like to have a model such that

$$\hat{f}_N(x) - \phi(x) \simeq 0,$$

i.e., a model with small **bias**

- This should be achievable if we have rich, highly flexible models
  - Or with essentially no regularization (or both!)
- But we would also like to have a robust model building procedure which results in a small **variance**  $V_N(x) \simeq 0$ 
  - This should be achievable if models are simple with few parameters
  - Or with more severe regularization if not
- Obviously both goals are contradictory to a large extent

- There is thus a **tradeoff** between bias (low for complex models) and variance (low for simple models)



Taken from *Hastie et al.*, p. 38

- Recall that in  $k$ -NN regression we choose the  $k$  nearest neighbors  $x^{p_1}, \dots, x^{p_k}$  of  $x$  and estimate  $\hat{y} = \hat{y}(x)$  as

$$\hat{y}(x) = \frac{1}{k} \sum_{j=1}^k y^{p_j}$$

- Here the parameter that controls the tradeoff is just  $k$
- If  $k = N$ , the sample size, the  $N$ -NN estimator is just the mean:  $Y_N^{NN} = \bar{y}$ , with very small variance but large bias (it's an obviously bad model!!)
- If  $k = 1$ , the 1-NN estimator will have smaller bias but a large variance: changing the sample is very likely to change the sample point nearest to  $x$

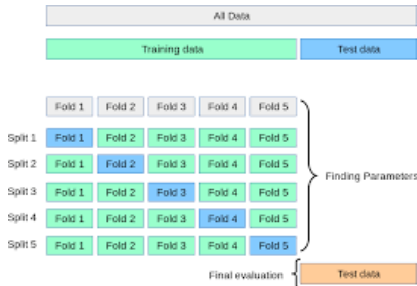
- Recall that the Ridge Regression loss is

$$e_R(w) = \frac{1}{2N} \sum_p (y^p - w \cdot x_p^p)^2 + \frac{\alpha}{2} \|w\|^2,$$

- Here the parameter that controls the tradeoff is  $\alpha$
- If  $\alpha \gg 1$ , any non zero  $w$  implies a large regularization penalty
  - It is thus likely that  $w \simeq 0$  and the Ridge model reduces again to the mean  $\bar{y}$ , with small variance but possibly a large bias
- But if  $\alpha \simeq 0$ ,  $w$  can wander on the entire  $\mathbf{R}^d$ 
  - The bias will be then smaller, but the weights  $w_S$  and  $w_{S'}$  from different samples may be very different, particularly if the covariance is nearly singular

- Recall that we build a model over a **train** sample but apply it on a **test** sample
- Thus, before we start applying a model, we should have a reasonably accurate idea of its performance in practice
- I.e., we want to estimate the model's **generalization performance**
- Estimating the generalization performance **only over the sample  $S$  used for training results in misleading error values**
- The preceding suggests to try to have  $M$  independent subsamples  $S_m$  and then
  - To compute  $\hat{f}_M(x) = \frac{1}{M} \sum_m \hat{f}_{S_m}(x) \simeq \hat{f}_N(x)$
  - To get the error estimate  $\hat{e} = \frac{1}{N} \sum_p (y^p - \hat{f}_M(x^p))^2$  over a new, **unseen** sample  $S' = \{(x^p, y^p)\}$
- But since usually we only have a single  $S$ , we “simulate” this by **Cross Validation** (CV) to get our first realistic generalization error estimates

- Graphically, the general schema of Cross Validation is the following



From [Scikit-learn](#)

- In more detail, in Cross Validation (CV) and for a global sample  $S$  we
  - Randomly split  $S = S_{tr-val} \cup S_{ts}$
  - Randomly split  $S_{tr-val}$  in  $M$  subsets  $S_1, \dots, S_M$
  - Work with  $M$  **folds**: pairs  $(S_m, S_m^c)$ , with

$$S_m^c = S - S_m = \bigcup_{i \neq m} S_i$$

- Build  $M$  different models **using the  $S_m^c$  as training subsets**
  - Compute their errors  $e_m$  on the folds' **validation subsets  $S_m$**
  - Use these errors' average as a first estimate of the true model performance
- CV can and **must be used** in any model building procedure
  - Most data science packages have tools to simplify this
- We will also use CV to find **optimal model hyper-parameters** such as  $\alpha$  in Ridge Regression



- Consider for Ridge regression a hyperparameter range  $\alpha \in [a, A]$ , with  $0 < a < A$ 
  - In general  $a \simeq 0$ ,  $A \gg 1$
  - $\alpha = a$ : essentially no penalty, small bias and high variance
  - $\alpha = A$ : large penalty, small variance but high bias
- Select an  $L + 1$  point **grid**  $\{\alpha_0 = a < \alpha_1 < \dots < \alpha_L = A\}$ 
  - The  $\alpha_j$  can be equi-spaced, log equi-spaced, random, ...
- At each  $\alpha_\ell$ 
  - Train  $M$  models on the  $S_m^c$  using the hyperparameter  $\alpha_\ell$
  - Average their test errors  $e_m$  on the  $S_m$  to get the error  $e(\alpha_\ell)$  at  $\alpha_\ell$
- And choose the final (hopefully) optimal hyperparameter  $\alpha^*$  as

$$\alpha^* = \arg \min_{0 \leq \ell \leq L} e(\alpha_\ell)$$

- $\alpha^*$  gives the model with the **best expected generalization among all possible  $\alpha$  choices**

- 1 We have stressed that **any model estimation is sample-dependent** and that this has to be controlled
- 2 We have introduced the **bias** and **variance** as the two key components of any model error
- 3 We have discussed **bias-variance trade-off**
- 4 We have introduced **Cross Validation** here as a tool to estimate a **model's generalization performance**
- 5 We have also introduced **Cross Validation** as a tool to estimate a **model's hyper-parameters**

## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- **Basic Classification**
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

- Recall that in regression we have numerical continuous targets  $y$  and want our predictions  $\hat{y}$  to be as close to  $y$  as possible
- But in classification we have a finite number of labelled targets for which “selection by closeness” doesn’t make sense
- Natural alternative: select the **most probable** label given the pattern  $x$  we have just received
  - The concrete labels used for targets do not matter much anymore
  - Model learning should thus be “target” agnostic
  - And good probability estimates should be quite useful
- Let’s analyze this in an example

- We want to diagnose whether a person may have diabetes from some clinical measures
- Features  $x$ : clinical measures
  - `numPregnant`
  - `bloodPress`
  - `massIndex`
  - `age ...`
- Target  $y$ : 0 (no diabetes), 1 (diabetes)
- Clear goal but perhaps too radical
- Better: try to estimate the probability  $P(1|x)$  of having diabetes depending on the features  $x$  we measure

- We have random patterns  $\omega$  from  $M$  classes,  $C_1, \dots, C_M$
- Over each pattern we “measure”  $d$  features  $x = x(\omega) \in \mathbb{R}^d$ 
  - $x$  inherits the randomness in  $\omega$  and becomes a random variable
- A  $\omega$  has a **prior probability**  $\pi_m$  of belonging to  $C_m$
- Inside each class  $C_m$  there is a **conditional class density**  $f(x|m)$  that “controls” the appearance of a given  $x$
- The  $\pi_m$  and  $f(x|m)$  determine the **posterior probability**  $P(m|x)$  that  $x$  comes from class  $C_m$
- **Intuition:** we should assign  $x$  to the class with the largest  $P(m|x)$ , that is, work with the classifier

$$\delta(x) = \arg \max_m P(m|x)$$

- **Bayes rule:**  $P(B|A) = \frac{P(A \cap B)}{P(A)}$
- This requires to work with probabilities, not densities, but  $P(\{x\}) = P(m \cap \{x\}) = 0$  and

$$P(m|x) = \frac{P(m \cap \{x\})}{P(\{x\})} = \frac{0}{0} = \dots???$$

- But we can use the approximation

$$\begin{aligned} P(m|x) &\simeq P(m|B_r(x)) = \frac{P(C_m \cap B_r(x))}{P(B_r(x))} = \frac{P(B_r(x)|m)P(C_m)}{P(B_r(x))} \\ &= \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))} = \pi_m \frac{\int_{B_r(x)} f(y|m) dy}{\int_{B_r(x)} f(z) dz} \end{aligned}$$

where we assume that features  $x$  are measured independently from classes  $m$

- Remember the Fundamental Theorem of Calculus:

if  $F(x) = \int_a^x f(y)dy$ ,

$$\lim_{\epsilon \rightarrow 0} \frac{1}{2\epsilon} \int_{x_0-\epsilon}^{x_0+\epsilon} f(y)dy = \frac{dF}{dx}(x_0) = f(x_0)$$

- In  $d$  dimensions it becomes

$$g(w) = \lim_{r \rightarrow 0} \frac{1}{|B_r(w)|} \int_{B_r(w)} g(z)dz$$

- Putting everything together, we arrive

$$\begin{aligned} P(m|x) &= \lim_{r \rightarrow 0} P(m|B_r(x)) = \pi_m \lim_{r \rightarrow 0} \frac{\int_{B_r(x)} f(y|m)dy}{\int_{B_r(x)} f(z)dz} \\ &= \pi_m \lim_{r \rightarrow 0} \frac{\frac{1}{|B_r(x)|} \int_{B_r(x)} f(y|m)dy}{\frac{1}{|B_r(x)|} \int_{B_r(x)} f(z)dz} = \frac{\pi_m f(x|m)}{f(x)} \end{aligned}$$



- This suggests that a natural choice is the **classifier**  $\delta_B$

$$\begin{aligned}\delta_B(x) &= \arg \max_m P(m|x) = \arg \max_m \frac{\pi_m f(x|m)}{f(x)} \\ &= \arg \max_m \pi_m f(x|m)\end{aligned}$$

- With some extra work we can show that this **Bayes Classifier**  $\delta_B$  defines an optimal solution (in some precise sense) of the classification problem
- But ... this doesn't look too practical, for we do not know either  $\pi_m$  or (much harder)  $f(x|m)$

- To define  $\delta_B$  we need to know the prior probabilities  $\pi_m$  and the prior densities  $f(x|m)$
- A reasonable choice for  $\pi_m$  is  $\hat{\pi}_m = \frac{N_m}{N}$ , where  $N_m$  is the number of patterns of  $C_m$  in the sample
- But effective multidimensional density estimates are rather difficult, because of the **curse of dimensionality**
  - Densities generalize histograms
  - Good histograms need accurate counts of elements nearby
  - But in high dimensions there won't be nearby elements!!
- Options:
  - Restrict possible density models: **logistic regression**
  - Assume no model and apply a **Nearest Neighbor (NN) strategy**

- Very simple: at any  $x$  consider the subset  $N_k(x)$  of its  $k$  closest sample points and
  - Let  $n_m(x)$  the number of elements of class  $m$  in  $N_k(x)$
  - Notice that  $0 \leq n_m(x) \leq k$
  - Define  $\delta_{kNN}(x) = \arg \max_m n_m(x)$
- That is,  $\delta_{kNN}(x)$  assigns  $x$  to the class that has more patterns in  $N_k(x)$
- We can partially justify this definition from a Bayesian point of view
- Assume that  $B_r(x)$  is the smallest ball that contains  $N_k(x)$  and consider the approximations
  - $P(B_r(x)|m) \simeq \frac{n_m(x)}{N_m}$
  - Similarly,  $P(B_r(x)) \simeq \frac{k}{N}$
  - And  $\pi_m \simeq \frac{N_m}{N}$

- We then have

$$\begin{aligned} P(m|x) &\simeq P(m|B_r(x)) = \frac{\pi_m P(B_r(x)|m)}{P(B_r(x))} \\ &\simeq \frac{N_m}{N} \frac{n_m(x)}{N_m} \frac{1}{\frac{k}{N}} = \frac{n_m(x)}{k} \end{aligned}$$

- Therefore  $\delta_{kNN}$  might be close to  $\delta_B$ , for

$$\begin{aligned} \delta_{kNN}(x) &= \arg \max_m n_m(x) = \arg \max_m \frac{n_m(x)}{k} \\ &\simeq \arg \max_m P(m|x) = \delta_B(x) \end{aligned}$$

- **Q1: How do we choose  $k$ ?** Using CV, of course
- There are no closed form solution and we have to deal again with the bias–variance tradeoff
  - Small variance with large  $k$ : if  $k = N$ ,  $k$ -NN classification returns the majority class
  - Small bias with small  $k$ : if  $k = 1$  a point very close to  $x$  should be in the same class
  - But also large variance: the nearest point to  $x$  in another sample may well belong to a different class
- **Q2: Is  $k$ -NN always meaningful?**
  - Well, only if classes are concentrated and far from each other
  - In other words, when the classification problem is easy!!

- Consider a two class problem with labels  $y = 0, 1$
- We will call patterns with label 1 **positive** and those with label 0 **negative**
  - Usually the positive patterns are the interesting ones: sick people, defaulted loans, ...
- Let  $\hat{y} = \hat{y}(x)$  the label predicted at  $x$ ; we say that  $x$  is a
  - **True Positive (TP)** if  $y = \hat{y} = 1$
  - **True Negative (TN)** if  $y = \hat{y} = 0$
  - **False Positive (FP)** if  $y = 0$  but  $\hat{y} = 1$
  - **False Negative (FN)** if  $y = 1$  but  $\hat{y} = 0$
- These values are combined in many performance measures

- The classifier **accuracy** is  $acc = \frac{TP+TN}{N}$ 
  - The first thing to measure but ...
- It may not be too significant when the number  $N_0$  of negatives is  $\gg N_1$ , the number of positives; then
  - The classifier  $\delta(x) = 0$  will have a high accuracy  $N_0/N \simeq 1$
  - But it will also be useless!!
- First variant: Precision, Recall
  - **Recall:**  $TP/(TP + FN)$ , i.e., the fraction of positives detected
  - **Precision:**  $TP/(TP + FP)$ , i.e., the fraction of true alarms issued
- Recall measures how many positive cases we recover, i.e., how effective is our method
- Precision measures the effort we need for that, i.e., its efficiency
- Ideally: high recall, high precision (i.e., effective and efficient!!)
- Another way of presenting the performance is through the **confusion matrix**

## ■ Standard layout

	P' (Predicted)	N' (Predicted)
P (Actual)	True Positive	False Negative
N (Actual)	False Positive	True Negative

## ■ Other layouts:

- **Positives (with label 1) at bottom** (as done in `confusion_matrix` of `sklearn`)
  - Predicted values in rows, real values in columns
- Accuracy and the confusion matrix can be extended to multi-class problems
  - But not other concepts such as recall or precision



## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- **Logistic Regression**

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

- $k$ -NN Classifier is simple but also crude; have to look elsewhere
- Building a regression model with targets given by some coding of class labels usually doesn't make sense
- However, for a binary 0–1 response, it can be shown that the  $w_0 + w \cdot x$  obtained using linear regression is in fact an estimate of  $P(1|x)$ 
  - We may thus fix a threshold  $\delta_0$  and decide 0 if  $w_0 + w \cdot x < \delta_0$  and 1 otherwise
  - However, we may end up with probability estimates less than 0 or bigger than 1!!!
- Better idea: try to transform the linear output  $w_0 + w \cdot x \in (-\infty, \infty)$  into a probability  $P(1, x) \in (0, 1)$ 
  - Logistic regression does this through the **logistic** or **sigmoid** function

- We assume

$$P(1|x) = P(1|x; w_0, w) = \frac{1}{1 + e^{-(w_0 + w \cdot x)}}$$

- Then  $0 \leq P(1|x) \leq 1$  for any  $x$
- We then have

$$P(0|x) = 1 - P(1|x) = \frac{e^{-(w_0 + w \cdot x)}}{1 + e^{-(w_0 + w \cdot x)}} = \frac{1}{1 + e^{w_0 + w \cdot x}}$$

- Notice that if  $w_0 + w \cdot x = 0$ ,  $P(1|x) = P(0|x) = 0.5$
- The ratio  $\frac{P(1|x)}{P(0|x)} = e^{w_0 + w \cdot x}$  is called the **odds** of  $x$  and its log the **log odds** or **logit**
- Thus, the basic assumption in LogR is that the **logit is a linear function**  $w_0 + w \cdot x$  of  $x$
- We have the model  $f(x; w)$ ; we need a **loss** function  $L(w_0, w)$  to minimize for which we use the sample's **likelihood**

- Assume a single sample  $x, y$  and two possible model coefficients  $w_0, w$  and  $w'_0, w'$
- Denoting by  $p = P(y|x; w_0, w)$  and  $p' = P(y|x; w'_0, w')$ , it is clear that we should prefer  $w_0, w$  if  $p > p'$  and  $w'_0, w'$  if not
  - In other words, we prefer the coefficients that give a **higher posterior probability**
- For an **independent** sample  $S = \{(x^p, y^p)\}$ , its joint probability under a posterior model  $p = P(y|x, w_0, w)$  is

$$P(Y|X; w_0, w) = \prod_{p=1}^N P(y^p|x^p; w_0, w)$$

- And, again, given two possible model coefficients  $w_0, w$  and  $w'_0, w'$ , **we should prefer  $w_0, w$  iff**

$$P(Y|X; w_0, w) > P(Y|X; w'_0, w')$$

- Therefore, we can estimate the optimal  $w_0^*$ ,  $w^*$  as

$$w_0^*, w^* = \arg \max_{w_0, w} P(Y|X; w_0, w)$$

- By the independence assumption we have

$$\begin{aligned} P(Y|X; w_0, w) &= \prod_{p=1}^N P(y^p|x^p; w_0, w) \\ &= \left\{ \prod_{y^p=1} P(1|x^p) \right\} \left\{ \prod_{y^p=0} P(0|x^p) \right\} \\ &= \prod_{p=1}^N P(1|x^p)^{y^p} P(0|x^p)^{1-y^p} \end{aligned}$$

because

- If  $y^p = 1$ ,  $P(1|x^p) = P(1|x^p)^{y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$ , and
- If  $y^p = 0$ ,  $P(0|x^p) = P(0|x^p)^{1-y^p} = P(1|x^p)^{y^p} P(0|x^p)^{1-y^p}$

- The **log–likelihood** of  $w_0, w$  given  $S$  is then

$$\begin{aligned}
 \ell(w_0, w; S) &= \log P(Y|X; w_0, w) \\
 &= \sum_p \{y^p \log p(1|x^p) + (1 - y^p) \log p(0|x^p)\} \\
 &= \sum_p y^p \log \frac{p(1|x^p)}{p(0|x^p)} + \sum_p \log p(0|x^p) \\
 &= \sum_p y^p (w_0 + w \cdot x^p) - \sum_p \log(1 + e^{w_0 + w \cdot x^p})
 \end{aligned}$$

- We can thus estimate the optimal  $\hat{w}_0^*, \hat{w}^*$  as

$$\hat{w}_0^*, \hat{w}^* = \arg \min_{w_0, w} -\ell(w_0, w; S)$$

- Extra bonus:  $-\ell$  is a convex differentiable function of  $(w_0, w)$  and, thus, it is enough to solve  $\nabla \ell(w_0, w) = 0$

- However,  $\nabla \ell(w_0, w) = 0$  doesn't admit a closed form solution, but only an iterative, numerical one
- We apply the **Newton–Raphson** iterative method, here equivalent to the general Newton method for function minimization
- For a general weight vector  $w$ , and starting with an initial random  $w^0$ , Newton's iterations are

$$w^{k+1} = w^k + (\mathcal{H}_\ell(w^k))^{-1} \nabla \ell(w^k)$$

- $\mathcal{H}_\ell(w^k)$  denotes the **Hessian** of  $\ell$  at  $w^k$ , which may or may not be invertible
  - Everything is fine if the  $w^k$  are close enough to the optimum  $w^*$  but far away things may get tricky
- Just as before, we can add a regularization term  $\frac{\alpha}{2} \|w\|^2$  to avoid invertibility problems
- The iterations in Logistic Regression are again typical of many of the model building methods used in Machine Learning

- The general approach to **learning** is the following:
  - A **model**  $f(x; w)$  is chosen
  - Given a sample  $S = \{(x^1, y^1), \dots, (x^N, y^N)\}$ , we define a **sample dependent loss function**

$$L(w) = L(w|S) = L(y^1, \dots, y^N, f(x^1; w), \dots, f(x^N; w))$$

- $L(w)$  is often minimized from some  $w^0$  by **iterations**

$$w^{k+1} = w^k - \rho_k G(w^k, S)$$

with  $\rho_k$  a **learning rate** and  $G$  some vectorial function

- When  $G(w) = \nabla L(w)$  we have **gradient descent**
  - When  $G(w) = \mathcal{H}(w)^{-1} \nabla L(w)$  we obtain **Newton's method**
- When the entire sample  $S$  is used at each iteration, we speak of **batch learning**
- When only single patterns  $(x^p, y^p)$  or small subsamples are used, we speak of **on-line** or **minibatch learning**



- We have just built our second neural network: the logistic regression architecture can be seen as a multilayer perceptron (MLP) with no hidden layers and sigmoid outputs
- Again, thinking of this as an MLP seems far-fetched, as we are missing the hidden layers
- Now the gradient computations are more involved
- But we also have learning as iterative **loss** minimization, i.e., **training**
- In fact, we will see that classification MLPs for two class problems perform **logistic regression on the last hidden layer outputs**
- And , again, the process from the input to the last hidden layer produces an **enhanced feature representation** upon which the regression error should be better

- Some things change from regression, some don't
- We should check feature correlations: although less clearly influential, they will affect most models
  - Scatter plots  $(x_i, y)$  are usually less informative
- Important: **positive and negative-class feature histograms**
- The **bias-variance trade-off** is subtler in classification
- Accuracy, recall, precision are the usual model quality measures
  - But NN training does not directly minimize them
- We use CV with **stratified folds** to estimate generalization performance
- We also use CV for hyperparameter estimation, as regularization will also be needed
  - In LogR we should minimize  $-\ell(w_0, w; S) + \frac{\alpha}{2} \|w\|^2$

- If possible, we don't want labels as model outputs but **posterior probabilities**
- Most models give them as pairs

$$(\hat{P}(0|x), \hat{P}(1|x)) = (\hat{P}(0|x), 1 - \hat{P}(0|x))$$

- In principle we would decide 1 if  $\hat{P}(1|x) > 0.5$  and viceversa, but this may be too crude
- It may be advisable to set a **decision threshold**  $0 < \theta < 1$  and decide 1 if  $\hat{P}(1|x) > 1 - \theta$  and 0 if  $\hat{P}(1|x) < \theta$
- For **imbalanced** problems where  $\pi_0 \gg \pi_1$  (usually the interesting ones) we would have  $\hat{P}(1|x) \simeq 0$  for most  $x$ 
  - In this case we may choose another  $\theta \ll 0.5$  and **suggest** 1 if  $\hat{P}(1|x) > \theta$

- 1 We have introduced the classification problem as one of **computing posterior probabilities**
- 2 We have found the **optimal Bayes classifier** and **approximated it by  $k$ -NN**
- 3 We have introduced **Logistic Regression** and the numerical minimization of its (minus) log-likelihood
- 4 We have introduced several measures of **classifier performance**
- 5 We have reviewed some **practical issues** of classification


## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons


- **Rosenblatt's Perceptrons**
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

## ■ Basic model: Ramón y Cajal's neuron (1900)




- Basic behavior: the neuron either fires or stays at rest depending basically on its inputs
- This translates into an electrical nerve impulse or **action potential**
- The brain has about  $10^{11}$  neurons
  - Each one has about 7,000 connections
  - These connections are often **recurrent**

- They developed (circa 1935) the first electro-physiological model to describe the generation and propagation of action potentials in neurons
- They used the giant squid axon for this



From Wikipedia's [Squid Giant Synapse](#)

## ■ Electronic version: McCulloch–Pitts (1940)



- Since  $I = \frac{V}{R} = wV$ , with  $w$  the conductance, the McCulloch–Pitts neuron output is

$$H \left( \sum_{j=1}^d w_j V_j + I \right) = V$$

with the Heaviside function  $H$  ensuring a 0–1 output

- The **perceptron** is a mathematical formulation of the McC–P neuron where **the conductances are to be learned**



- Given a sample  $S = \{(x^p, y^p)\}$  with  $y^p = \pm 1$ , **Rosenblatt's Perceptrons** (PCPs) are linear machines  $w \cdot x$  such that

$$\begin{aligned} w \cdot x^p &> 0 & \text{if } y^p &= 1 \\ w \cdot x^p &< 0 & \text{if } y^p &= -1 \end{aligned}$$

- This can be seen as a **homogeneous** classification problem
- In more compact form we want for all  $p$

$$y^p w \cdot x^p > 0$$

- If such a **separating**  $w$  exists, it can be computed in many ways
- But Rosenblatt's goal was to **learn**  $w$ , i.e., to arrive at a separating  $w$  by repeatedly examining the  $(x^p, y^p)$  and adjusting  $w$  if necessary

## ■ The **Delta Rule** algorithm

Start with  $w = 0$

While stopping condition not met:

    Get a new pattern  $(x^p, y^p) = (x^{p(t)}, y^{p(t)})$

    if  $y^p w \cdot x^p \leq 0$  then:

$w = w + y^p x^p$

- Writing  $w$  for the new vector and  $w'$  for the previous one, notice that

$$y^p w \cdot x^p = y^p w' \cdot x^p + \|x^p\|^2 > y^p w' \cdot x^p$$


- It states that if the sample  $S$  is **linearly separable**, the Delta rule yields a separating  $w$  in a **finite number of steps**
  - I.e., it **learns!!**
- More precisely, if for some separating hyperplane  $w$  with  $\|w\| = 1$  we have  $\min_p y^p w \cdot x^p = \gamma > 0$ , then **a separating  $w'$  can be found after at most  $T = R^2/\gamma^2$  wrong classifications, where  $R = \max_p \|x^p\|$**
- Notice that  $\gamma$  is the **margin** of the  $w$  hyperplane. Thus, if

$$\gamma^* = \max_{\|w\|=1} m(w, S) = \max_{\|w\|=1} \min_p y^p w \cdot x^p$$


is the **maximum margin**, we will need at most  $R^2/(\gamma^*)^2$  iterations to train a Perceptron

- The maximum margin hyperplane gives the best (smallest) Novikov bound
- And problems with smaller margins require more iterations

- Intuitively, when the margin is large, the problem is easier



- The margin of  $w$  is just the smallest distance between sample points and the hyperplane defined by  $w$



- This extends to the multidimensional case: if  $\|w\| = 1$  and  $b = 0$ , the distance of  $x$  to the hyperplane defined by  $w$  is  $|w \cdot x| = yw \cdot x$
- Thus the Delta Rule **corrects weight vectors with negative margins**

- Assume  $x^t$  is the  $t$ -th pattern such that  $y^t w^{t-1} \cdot x^t \leq 0$
- Then  $\|w^t\|^2 = \|y^t x^t + w^{t-1}\|^2 \leq \|x^t\|^2 + \|w^{t-1}\|^2$  and, thus,

$$\|w^t\|^2 \leq \|x^t\|^2 + \dots + \|x^1\|^2 + \|w^0\|^2 \leq tR^2$$

if  $w^0 = 0$  and  $\|x^p\| \leq R$  for all  $p$

- Also, if  $w'$  is a 1-norm separating hyperplane i.e.,  $\min_p y^p w' \cdot x^p = \gamma > 0$  for all  $p$ , we have

$$w^t \cdot w' = (y^t x^t + w^{t-1}) \cdot w' \geq \gamma + w^{t-1} \cdot w' \geq \dots \geq t\gamma + w^0 \cdot w'$$


- Putting both together we have for all  $t$  for which  $x^t$  is erroneous

$$t\gamma \leq w^t \cdot w' \leq \|w^t\| \leq \sqrt{t}R$$


which implies  $t \leq R^2/\gamma^2$

- Thus, the Delta Rule will stop after at most  $R^2/\gamma^2$  iterations

- Rosenblatt's Perceptron caused a big stir on the feasibility of **Conexionist Artificial Intelligence** but soon its first winter came
- Observation 1: A PCP can learn (obviously) only linearly separable problems; for instance, it can learn the AND and OR predicates



- But a **PCP** cannot learn the **XOR** predicate




- Observation 2: for PCPs to be useful, **linearly separable problems must be frequent**
- Q: how frequent are they?
- A **dichotomy** for  $S = \{x^1, x^2, \dots, x^N\}$  is any separation of  $S$  in two classes
- It corresponds to any assignment of the  $\pm 1$  values to the  $x^p$
- The total number of dichotomies is  $2^N$
- Thus, the **total number of 2-class problems over  $S$**  is  $2^N$




- We have to compare the total number of dichotomies with the number of dichotomies that are **linearly separable**
- The points in a sample  $S$  are said to be in **general position** if there are not  $d + 1$  patterns in a  $d - 1$ -dimensional hyperplane
- **Theorem** If  $S$  is in general position, the number  $L(N, d)$  of linearly separable dichotomies is

$$L(N, d) = \left\{ \begin{array}{ll} 2^N & \text{si } N \leq d + 1 \\ 2 \sum_{i=0}^d \binom{N-1}{i} & \text{si } N \geq d + 1 \end{array} \right\}$$

- Consider  $d = 2$ ,  $3 = d + 1$  points and a  $1 = d - 1$ -dimensional hyperplane



- Consider now  $d = 2$  and  $3 = d + 1$  points **not** on a  $1 = d - 1$ -dimensional hyperplane (i.e., a line)



- Notice that for  $d$  fixed,

$$\frac{L(N, d)}{2^N} \rightarrow 0$$

as  $N \rightarrow \infty$

- In practice we can expect  $N \gg d$  and the fraction of separable dichotomies will be very small
- Thus, linear PCPs will not be useful
- But a **one hidden layer** PCP can solve XOR and, in fact, **separate any convex region** from the rest of the space
- And a **two hidden layer** PCP can separate any **polyhedral region** and hence **solve any classification problem**
- The same is essentially true for regression problems
- But: **no algorithm is known to learn these 2-hidden layer PCPs**

- Rosenblatt's Perceptrons are natural **online** machines:
  - Patterns appear sequentially and each may change the current weights
- We can also define a perceptron loss over a given sample as

$$L(w|S) = - \sum_{\{p: y^p w \cdot x^p < 0\}} y^p w \cdot x^p$$

- The local error would thus be  $-y^p w \cdot x^p$  when  $y^p w \cdot x^p < 0$  and 0 otherwise
- Since  $\nabla_w (-y^p w \cdot x^p) = -y^p x^p$ , the Delta Rule can be seen as a kind of local gradient descent on the loss  $L(w|S)$ :

$$w^t = w^{t-1} + y^{p(t)} x^{p(t)} = w^{t-1} - \nabla_w \left( -y^{p(t)} x^{p(t)} \cdot w \right)$$

- However  $L(w|S)$  is not differentiable and it is not clear what to make of  $\nabla_w L(w|S)$  ...

- First option: work with linear models but somehow ensure that  $d \gg N$ 
  - Transform the original  $d$ -dimensional features  $x$  in new  $D$ -dimensional ones such that  $D \gg d$
  - We overcome then that, in practice, we usually have  $N \gg d$
  - This is the goal of **kernel-based Support Vector Machines**
- Second option: build trainable PCPs replacing the Heaviside function
  - Work with differentiable PCP transfer function
  - Transform PCP training into a differentiable optimization problem
  - This is what we will do next
- Before doing so, note that **Rosenblatt's PCPs can be seen as the first Machine Learning model**
- And modern ML started around 1985 along the PCP's ideas


## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- **Classical MLPs**
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

- General layout:
  - An input layer (input)
  - One or several hidden layers
  - One output layer
- Feedforward connections only



- Example: [TensorFlow Playground](#)



- **No feedback or lateral connections**
- **Fully connected layers**
- **Linear weight connections** between units and (usually) **non linear activations** inside each unit
  - Combined effect of sucesive layers: potentially highly non-linear transformation
- General processing: layered and feedforward
- In practice (1990s), one hidden layer and only sometimes two
- Later (around 2010): Deep Networks with “many” (from 3 to 10) layers
- Now: almost anything!

- The **activations**  $a_i^h$  of a unit in layer  $h$  receives the **outputs**  $o_j^{h-1}$  from processing in the previous layer

$$a_i^h = \sum_{j=1}^{n_{h-1}} w_{ij}^h o_j^{h-1} + b_i^h,$$

- In matrix/vector form:

$$\mathbf{a}^h = \mathbf{W}^h \mathbf{o}^{h-1} + \mathbf{b}^h$$

- **Output**  $o_i^h$  of a unit: **non linear processing**  $o_i^h = \varphi(a_i^h)$  of its activation
- In vector form:

$$\mathbf{o}^h = \varphi(\mathbf{a}^h),$$

where  $\varphi$  is applied over each unit

## ■ Choices for $f$ :

- Heaviside (in the very first Rosenblatt's Perceptrons):  $\varphi(a) = 0$  if  $a \leq 0$ ,  $\varphi(a) = 1$  if  $a > 0$

- **Identity/linear**:  $\varphi(a) = a$

- **Sigmoid**:


$$\varphi(a) = \sigma(a) = \frac{1}{1 + e^{-a}}$$

- **Hyperbolic tangent**:

$$\varphi(a) = \tanh(a) = \frac{e^a - e^{-a}}{e^a + e^{-a}}$$

- **Rectified Linear Units (ReLU)**:  $\varphi(a) = r(a) = \max(0, a)$

- Sigmoid and tanh: smooth version of Heaviside step function
- For ReLUs  $r'(x)$  is either 0 or 1 (hoping  $x = 0$  never happens!!)
  - Many gradient elements will go to 0 and many units will have constant activations no matter their input



From [Stanford's UFLDL Tutorial](#)

- Usual activation choices:
  - For hidden units: hyperbolic tangent and, currently, ReLU
  - Linear outputs for regression, sigmoid/soft max outputs for classification

- We need first to compute an MLP's outputs through the **forward pass**
- At each layer  $h$  we compute two variable sets
  - The linear **activations**  $a^h = W^h o^{h-1} + b^h$  with  $o^{h-1}$  the previous layer's output and  $W^h, b^h$  the **weights** and **biases** connecting layer  $h - 1$  with layer  $h$
  - The nonlinear **outputs**  $o^h = \varphi(a^h)$ ,
- In unit-wise terms:  $a_i^h = \sum w_{ij}^h o_j^{h-1} + b_i^h$ ,  $o_i^h = \varphi(a_i^h)$
- Concatenating these operations we arrive at the outputs  $\hat{y} = f(x; \mathcal{W})$  of the network, with  $\mathcal{W} = (W^1, b^1, \dots, W^H, b^H)$  the network's weight set

- The Single Hidden Layer (SHL) MLP
  - $D$  inputs (determined by the problem at hand)
  - One hidden layer with  $H$  units (number to be chosen) and  $\tanh$  (or what you prefer) activation
  - One linear or sigmoid output (according to the problem at hand)
- Input–hidden processing: denoting inputs by  $x$  and the hidden unit output as  $o^H$ ,

$$o_h^H = \tanh \left( b_h^H + \sum_{j=1}^D w_{hj}^H x_j \right)$$

- In matrix/vector form:  $o^H = \tanh (W^H x + b^H)$

- Hidden–output processing: since targets are 1–dimensional, we have for the outputs  $\hat{y}$

$$\hat{y} = \sum_{h=0}^H w_h^O o_h^H + b^O,$$

- In vector form:  $\hat{y} = w^O \cdot o^H + b^O$
- That is, a **linear model** on the last hidden layer outputs  $o^H$
- Global process:

$$\hat{y} = f(x; w^O, W^H, b^O, b^H) = b^O + \sum_h w_h^O \tanh \left( b_h^H + \sum_j W_{hj}^H x_j \right)$$

- Or in matrix/vector form

$$\hat{y} = f(x; w^O, W^H, b^O, b^H) = b^O + w^O \cdot \tanh \left( b^H + W^H x \right)$$

- Or in Python `y_pred = np.tanh(x.dot(w_H.T) + b_H).dot(w_O.T) + b_O`

- MSE is the standard error function for regression MLPs

$$\begin{aligned} e(\mathcal{W}) &= \frac{1}{2} E_{x,y} [(y - f(x; \mathcal{W}))^2] = E_{x,y} [e^\ell(x, y; \mathcal{W})] \\ &= \int e^\ell(x, y; \mathcal{W}) p(x, y) dx dy \end{aligned}$$

with  $e^\ell(x, y; \mathcal{W})$  denotes the **local error**

$$e^\ell(x, y; \mathcal{W}) = \frac{1}{2} (y - \hat{y})^2 = \frac{1}{2} (y - f(x; \mathcal{W}))^2$$



- The general idea would be to obtain the optimal  $\mathcal{W}^*$  as a solution of  $\nabla e(\mathcal{W}) = 0$ , where

$$\begin{aligned}\nabla e(\mathcal{W}) &= E_{x,y} [\nabla_{\mathcal{W}} e^{\ell}(x, y; \mathcal{W})] \\ &= E_{x,y} [\nabla_{\mathcal{W}} f(x; \mathcal{W})(f(x; \mathcal{W}) - y)]\end{aligned}$$

for we have

$$\begin{aligned}\nabla_{\mathcal{W}} e^{\ell}(x, y; \mathcal{W}) &= -(y - f(x; \mathcal{W})) \nabla_{\mathcal{W}} f(x; \mathcal{W}) \\ &= \nabla_{\mathcal{W}} f(x; \mathcal{W})(f(x; \mathcal{W}) - y)\end{aligned}$$

- This suggests to get  $\nabla_{\mathcal{W}} f(x; \mathcal{W})$  but it is much better to focus on  $\nabla e^{\ell}$ , and
  - We will compute first  $\nabla e^{\ell}$  (and then  $\nabla e = E[\nabla e^{\ell}]$ )
  - We will exploit it to build MLPs through **optimization methods**

- Recall that we compute an MLP's outputs through the **forward pass**
- At each layer  $h$  we compute (and **store**) first the linear **activations**  $a^h = w^h o^{h-1} + b^h$  with  $o^{h-1}$  the previous layer's output and  $W^h, b^h$  the **weights** and **biases** connecting layer  $h - 1$  with layer  $h$
- We then compute (and store) the nonlinear **outputs**  $o^h = \varphi(a^h)$ ,
- In unit-wise terms:

$$a_i^h = \sum w_{ij}^h o_j^{h-1} + b_i^h, \quad o_i^h = \varphi(a_i^h)$$

- And we also have

$$\frac{\partial a_i^h}{\partial w_{ij}^h} = o_j^{h-1}, \quad \frac{\partial o_i^h}{\partial a_i^h} = \varphi'(a_i^h)$$

- The key tool is the **chain rule** that we will apply backwards from the output layer
- If  $w_{ij}^h$  is the general weight connecting unit  $j$  in layer  $h - 1$  to unit  $i$  in layer  $h$ , we have

$$\frac{\partial e^\ell}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} \frac{\partial a_i^h}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} o_j^{h-1} = \delta_i^h o_j^{h-1}$$

- That is, we have an easy to compute term  $\frac{\partial a_i^h}{\partial w_{ij}^h}$ , and a more complicated but crucial one  $\frac{\partial e^\ell}{\partial a_i^h}$
- We refer to  $\frac{\partial e^\ell}{\partial a_i^h} = \delta_i^h$  as the **generalized error** at unit  $i$  of layer  $h$

# Gradient at the Output Layer

Machine Learning Basics Multilayer Perceptrons

- Let's assume that we have layers  $1, 2, \dots, H-1, H$ , with  $H$  the output layer
- For simplicity, consider a regression problem with linear outputs
- In the output  $H$  layer we have

$$e^\ell = \frac{1}{2}(y - \hat{y})^2 \text{ and } \hat{y} = o^H = a^H = \sum_j w_j^H o^{H-1}$$

- Thus, the partial derivative is straightforward

$$\delta^H = \frac{\partial e^\ell}{\partial a^H} = \frac{\partial e^\ell}{\partial \hat{y}} = \hat{y} - y$$

i.e., the generalized error  $\delta^H$  is here the **standard error**

- Therefore, we have for the output layer  $H$

$$\frac{\partial e^\ell}{\partial w_j^H} = (\hat{y} - y) \frac{\partial a^H}{\partial w_j^H} = (\hat{y} - y) o_j^{H-1}$$

- Now it is (quite) easy to extend this to other layers

- Let  $w_{ij}^h$  be the weight connecting unit  $j$  in layer  $h - 1$  with unit  $i$  in layer  $h$ ; recall that then


$$\frac{\partial e^\ell}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} \frac{\partial a_i^h}{\partial w_{ij}^h} = \frac{\partial e^\ell}{\partial a_i^h} o_j^{h-1}$$

- And if we have already computed the generalized errors  $\delta_k^{h+1} = \frac{\partial e^\ell}{\partial a_k^{h+1}}$  for layer  $h + 1$ , we have

$$\begin{aligned} \frac{\partial e^\ell}{\partial a_i^h} &= \sum_k \frac{\partial e^\ell}{\partial a_k^{h+1}} \frac{\partial a_k^{h+1}}{\partial a_i^h} = \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial a_i^h} \\ &= \sum_k \delta_k^{h+1} \frac{\partial a_k^{h+1}}{\partial o_i^h} \frac{\partial o_i^h}{\partial a_i^h} = \left[ \sum_k \delta_k^{h+1} w_{ki}^{h+1} \right] \varphi'(a_i^h) \end{aligned}$$

- The previous formulae work for any number of hidden layers
  - We can work with **deep** MLPs
- Moreover, the chain rule can be automated and modern NN packages compute gradients automatically for very general NN architectures
  - Basic tool: automatic symbolic differentiation on computational graphs
  - This is one of the key reasons of the great success of deep neural nets
- Some examples: convolutional nets, residual nets, recurrent NNs, ...

- To deal with such general networks requires **automating backprop**
- This is done upon the network's **computational graph**



From Zhang et al., **Dive into Deep Learning**

- It corresponds to a SHL with a loss  $J = L + s$
- The circles contain differentiable operators
- The shadowed squares contain the variables against which differentiation takes place

## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- **Neural Network Classification**
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs



- Recall that our Logistic Regression model gave us estimate of  $P(0|x; w_0, w)$ ,  $P(1|x; w_0, w)$
- We then wrote the likelihood of a given sample as

$$P(Y|X; w_0, w) = \prod_{p=1}^N P(y^p|x^p; w_0, w)$$

- And then derived the optimum parameters  $w_0^*$ ,  $w^*$  as

$$w_0^*, w^* = \arg \max_{w_0, w} P(Y|X; w_0, w)$$

- Notice that we can see the output of a logistic regression model as that of an **NN with no hidden layers and sigmoid outputs**

- We can follow the same path for a 2 class classification problem changing that output for the one

$$f(x; \mathcal{W}) = w_0 + w \cdot \Phi(x, \mathcal{W}^H)$$

with  $w_0, w$  the weights from the last hidden layer (LHL) to the output and  $\Phi(x, \mathcal{W}^H)$  the map from inputs to the LHL outputs

- And now derive the optimum parameters  $\mathcal{W}^*$  as

$$\begin{aligned} \mathcal{W}^* &= \arg \min_{\mathcal{W}} - \sum_{p=1}^N \log P(y^p | x^p; \mathcal{W}) \\ &= \arg \min_{w_0, w, \mathcal{W}^H} - \sum_p y^p (w_0 + w \cdot \Phi(x^p, \mathcal{W}^H)) + \\ &\quad \sum_p \log(1 + e^{w_0 + w \cdot \Phi(x^p, \mathcal{W}^H)}) \end{aligned}$$

- Forward and backpropagation are also straightforward here
- We change things a little bit for  $K$  class problems

## Machine Learning Basics Multilayer Perceptrons

- We consider an input layer and a number of hidden layers
- Targets are the 1-hot encodings of class labels, so we use  $K$  outputs
- We want the MLP's  $k$ -th output to estimate the posterior  $P(k|x)$
- The natural output layer activation is thus the **softmax function**

$$\sigma_j(o) = F_j(x; \mathcal{W}) = \frac{e^{w_0^j + w^j \cdot o}}{\sum_1^K e^{w_0^k + w^k \cdot o}}$$

- Here  $o$  are the last hidden layer outputs and  $w^j$ ,  $w_0^j$  are the weight and bias from the last hidden layer to the network's  $j$  output
- Moreover  $\sum_j F_j(x; \mathcal{W}) = 1$  and we assume  $P(j|x) \simeq F_j(x; \mathcal{W})$
- For two classes this becomes

$$\sigma_1(o) = \frac{e^{w_0^1 + w^1 \cdot o}}{e^{w_0^0 + w^0 \cdot o} + e^{w_0^1 + w^1 \cdot o}} = \frac{1}{1 + e^{w_0^0 - w_0^1 + (w^0 - w^1) \cdot o}} = \frac{1}{1 + e^{\omega_0 + \omega \cdot o}}$$

- We thus get the **sigmoid activation** of Logistic Regression

- We change slightly the preceding for a  $K$  class classification problem
- We define an MLP  $F(x; \mathcal{W}) \in \mathbf{R}^K$  with softmax outputs
  - Recall that  $0 \leq F_k(x; \mathcal{W}) \leq 1$  and  $\sum_k F_k(x; \mathcal{W}) = 1$
  - And also we assume that  $P(k|x) \simeq F_k(x; \mathcal{W})$
- We label now the multiclass targets via the 1–hot encoding, with the target of class  $k$  being the vector  $e_k = (0, \dots, \underbrace{1}_k, \dots, 0)$
- Then, if  $x \in C_k$ , its label is  $y = (y_1, \dots, y_K)^t$  is  $e_k$ , with  $y_k = 1$ , and

$$P(k|x) = P(k|x)^1 = P(k|x)^{y_k} = \prod_{c=1}^K P(c|x)^{y_c} \simeq \prod_{c=1}^K F_c(x; \mathcal{W})^{y_c}$$

- We have a sample  $S = \{x^p, y^p\}$  with  $(x^p, y^p)$  in class  $c(p)$  and, hence,  $y^p = e_{c(p)}$
- We also consider a posterior model  $P(c|x, \mathcal{W}) = F_c(x, \mathcal{W})$ , the probability of getting  $S = (X, Y)$  is

$$\begin{aligned}
 P(Y|X; \mathcal{W}) &= \prod_1^N P(c(p)|x^p; \mathcal{W}) = \prod_1^N \prod_{c=1}^K P(c|x^p; \mathcal{W})^{y_c^p} \\
 &\simeq \prod_1^N \prod_{c=1}^K F_c(x; \mathcal{W})^{y_c^p}
 \end{aligned}$$

- Thus, we estimate  $\mathcal{W}$  by maximizing the sample's **likelihood**

$$\mathcal{P}(Y|X; \mathcal{W}) = \prod_{p=1}^N \prod_{c=1}^K F_c(x^p; \mathcal{W})^{y_c^p}$$

- As before, we will work with the sample's log-likelihood, i.e.

$$\begin{aligned}\ell(\mathcal{W}; S) &= \log \mathcal{P}(Y|X; \mathcal{W}) = \sum_{p=1}^N \sum_{c=1}^K \log \left( F_c(x^p; \mathcal{W})^{y_c^p} \right) \\ &= \sum_{p=1}^N \sum_{c=1}^K y_c^p \log F_c(x^p; \mathcal{W})\end{aligned}$$

- The **cross-entropy** loss is now simply the negative log-likelihood  $-\ell(\mathcal{W}; S)$ , i.e.,

$$L(\mathcal{W}) = -\ell(\mathcal{W}; S) = - \sum_{p=1}^N \sum_{c=1}^K y_c^p \log F_c(x^p; \mathcal{W})$$

- It is now straightforward to carry this into a NN setting

- 1 They have a layered structure with outputs computed in a **forward pass** using differentiable activations
- 2 Usual activations: ReLUs, sigmoid, tanh, linear
- 3 MSE is the usual regression cost; we use a variant of cross entropy in classification
- 4 The error function gradients are computed by **backpropagation** of generalized errors
- 5 Backprop is basically a very simple procedure than can be **largely automated**
- 6 Regression and classification problems fit very well in this setting

## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- **MLP Regularization**
- Unconstrained Smooth Optimization
- Computational Costs of MLPs



- In regression we assume  $y = \phi(x) + n$  and our goal is to get  $f \simeq \phi$


- We say that  $\mathcal{F} = \{f(x; \mathcal{W})\}$  is a **Universal Approximation Family** over a domain  $\mathcal{R}$  if

For any  $\epsilon > 0$  and any reasonable  $\phi$ , we can find an  $f(x; \mathcal{W}_{\phi, \epsilon})$  s.t.

$$e(\mathcal{W}_{\phi, \epsilon}) = \int (\phi(x) - f(x; \mathcal{W}_{\phi, \epsilon}))^2 p(x) dx \leq \epsilon$$

- Notice that Universal Approximation is just what we need in regression
- In fact a **Single Hidden Layer MLP with enough hidden units is an effective universal approximator**
- But it also implies the risk of **overfit** whose control is crucial

- Since MLPs are a UAF, they can also approximate the noise in the sample
  - Given  $S = \{(x^p, y^p)\}$  if we allow enough hidden units in a SHL MLP we can arrive to a  $w^*$  s.t.  $y^p = f(x^p; w^*)$
  - We get thus a sample error  $\hat{e}(w^*) = 0$  but possibly with a very high generalization error
- I.e., MLPs may have very small bias but possibly large variance



- Why is there overfitting?
  - Because we may end up having too many weights with respect to sample size
  - Because we allow these weights to explore the entire weight space
- We can avoid this wandering if we limit  $w$ 's growth, for which we add a **regularization** term  $g(\|w\|)$  to  $e(w)$  with  $g(r)$  increasing
- Working with  $e_R(w) = e(w) + g(\|w\|)$  we have to **balance** the minimization of  $e(w)$  and that of  $g(\|w\|)$
- This balanced learning results in better generalization

- The simplest regularization procedure adds a quadratic penalty to the square error  $e$

$$e_R(w) = e(w) + \frac{\lambda}{2} \|w\|^2,$$

with  $\lambda$  the **weight decay** factor

- Also known as Tikhonov's regularization (or Ridge Regression for linear models)
- The desired effect is to constrain the evolution of  $w$ :
  - In fact, the regularized loss is essentially the **Lagrangian of the constrained problem**

$$\min_w e(w) \text{ subject to } \|w\|^2 \leq \rho, \quad \rho > 0$$

- The gradient becomes  $\nabla e_R(w) = \nabla e(w) + \lambda w$  and gradient descent becomes

$$w^{k+1} = w^k - \rho_k (\nabla_w e(w^k) + \lambda w^k)$$

- Again, the correct choice of  $\lambda$  is crucial
- A small  $\lambda \ll 1$  results in a small regularization effect and overfitting risk appears
- A large  $\lambda \gg 1$  causes learning to forget about  $e(w)$  and the model will be essentially constant and will underfit
- Usually  $\lambda$  is chosen by using CV by exploring a discrete set of values  $\lambda_j$ ,
- The same can essentially be done for any other hyper-parameter:
  - Number of hidden layers and of hidden units?
  - Learning rate??
  - Minibatch size???
- But this may become unmanageable if we want to work with many (more than three?) hyperparameters

- Recall that for MLPs  $e(w)$  does not have a single minimum
- Moreover, the final MLP depends on the random initial  $w^0$
- And mini-batch training adds extra randomness to the final model
- Therefore, **you may never get the same MLP twice!!!**
- But we can turn this to our advantage by
  - Starting from  $K$  independent initial weights and get  $K$  optimal weight sets  $w_k^*$
  - Giving as the output the average  $f_e(x) = \frac{1}{K} \sum_1^K f(x; w_k^*)$
- We expect outputs of the form  $\hat{y}_k^p = y^p + \epsilon_k^p$  with the errors  $\epsilon_k^p$  being independent
- Hence  $\frac{1}{K} \sum_k \epsilon_k^p \simeq 0$  and  $\frac{1}{K} \sum_k \hat{y}_k^p \simeq y^p$

- MLPs have a **high risk of overfitting**
- Thus, they must be **regularized** to avoid overfitting
  - The regularization hyperparameter is chosen through **cross validation**
- MLP training has two **random components**: the initial point and minibatch selection
  - Thus, each training will converge to a **different optimum**
- MLP **ensembles** can take advantage of this
  - They imply extra costs but ensembles are **embarrassingly parallelizable**

## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression


## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- **Unconstrained Smooth Optimization**
- Computational Costs of MLPs



- To build a ML model  $\equiv$  to **minimize a loss function**
- General optimization theory is a key tool in Machine Learning (ML)
- There are two optimization set ups in ML
  - **Unconstrained** optimization, slightly simpler and the one used for MLPs
  - **Constrained** optimization, wider and more complex
- In ML we have also to consider the optimization of differentiable and also non differentiable error functions
- MLP optimization: **unconstrained and differentiable**
- And also **batch**, i.e., over the entire sample, **mini-batch** over subsamples or **on line**, pattern by pattern

- We “stretch” the weight set  $\mathcal{W}$  into a vector  $w$
- We recall that  $-\nabla e(w)$  is the maximum descent direction
- First idea: to build a (hopefully convergent) sequence  $w^k$  iterating (small) steps along  $-\nabla e(w^k)$



- In more detail, we start from a random  $w^0$  and compute

$$w^{k+1} = w^k - \rho_k \nabla_w e(w^k)$$

- $\rho_k$  is the **learning rate** (LR)
- With a small  $\rho_k$  we ensure  $e(w^{k+1}) < e(w^k)$  (although with possibly a very small descent)
- We can get a better iteration  $w^{k+1} = w^k - \rho_k^* \nabla_w e(w^k)$  using a  $\rho_k^*$  given by

$$\rho_k^* = \arg \min_{\rho} e(w^k - \rho \nabla_w e(w^k));$$

this is known as **line minimization**

- These GD methods are called **first order methods** in part because they only use  $\nabla e$

- Assume a quadratic function  $q(w) = aw^2 + bw + c$ , with  $a > 0$ , and a minimum at  $w^*$
- We can reach  $w^*$  from any  $w$  with a step  $\Delta w$  such that

$$0 = q'(w + \Delta w) = 2a(w + \Delta w) + b$$

- We have thus  $\Delta w = \frac{-b-2aw}{2a}$ , that is,

$$w^* = w - \frac{2aw + b}{2a} = w - \frac{1}{q''(w)} q'(w)$$

- This leads to **Newton's method**: minimize a general  $f$  iteratively using steps

$$w^{k+1} = w^k - \rho_k \frac{1}{f''(w^k)} f'(w^k)$$

with  $\rho_k$  a suitable learning rate

- Notice how the  $\frac{1}{f''(w_k)}$  acts as a **self adjusting learning rate**
- When the parabola is very sharp (i.e.,  $a \gg 1$ ), we will be close to the minimum and the descent can overstep it
- But then  $f''(w_k)$  will be big and, hence,  $\frac{1}{f''(w_k)}$  **moderates** the gradient step
- Conversely, when the parabola is very wide (i.e.,  $a \simeq f''(w_k) \ll 1$ ), the minimum will be far away and the gradient step small descent will
- But then  $f''(w_k)$  will be rather small and, hence,  $\frac{1}{f''(w_k)}$  **accelerates** the gradient step
- We can extend this to several dimensions

- For a  $d$  dimensional  $w$ , the Taylor expansion of  $e$  at an optimum  $w^*$  is

$$e(w) \approx e(w^*) + \frac{1}{2}(w - w^*)^t \cdot \mathcal{H}(w^*) \cdot (w - w^*)$$

- $\mathcal{H}(w^*)$  is the **Hessian** of  $e$  at  $w^*$  and  $\nabla e(w^*) = 0$
- It follows that  $\nabla e(w) \approx \mathcal{H}(w^*) \cdot (w - w^*)$  and, therefore,

$$w^* \approx w - \mathcal{H}(w^*)^{-1} \nabla_w e(w)$$

- This suggest to derive the  $w^k$  by

$$w^{k+1} = w^k - \rho_k \mathcal{H}(w^k)^{-1} \nabla_w e(w^k)$$

which is known as **Newton's Method** (NM)

- Theoretically NM converges very fast near  $w^*$ , but
  - Far from  $w^*$  convergence is not guaranteed
  - Moreover  $\mathcal{H}(w^k)$  may not be invertible
  - Besides, computing  $\mathcal{H}(w^k)$  is cumbersome and costly
- The **Gauss–Newton (GN)** approximation

$$\mathcal{H}(w) \simeq E[\nabla e(w) \nabla e(w)^T],$$

holds for any quadratic cost and simplifies the third problem

- The **Levenberg–Marquardt (LM)** method deals with the other two problems combining
  - Gradient descent “away” from  $w^*$
  - Gauss–Newton “near”  $w^*$

although “away” and “near” have to be properly addressed

- Assume  $e(w) = \frac{1}{2} \int (f(x; w) - y)^2 p(x, y) dx dy$ ; then

$$e'(w) = \int (f(x; w) - y) \frac{\partial f}{\partial w}(x, y) p(x, y) dx dy;$$

$$e''(w) = \int \left( \frac{\partial f}{\partial w} \right)^2 p(x, y) dx dy \\ + \int (f(x; w) - y) \frac{\partial^2 f}{\partial w^2}(x, y) p(x, y) dx dy$$

- Near a minimum  $w^*$  we may expect  $f(x; w) \simeq y$  and, therefore,

$$e''(w) \simeq \int \left( \frac{\partial f}{\partial w} \right)^2 p(x, y) dx dy > 0$$

- Thus, for square errors, we can use first derivatives to approximate  $e''(w)$



- In the general case we have

$$\nabla e(w) = E [\nabla f(x; w) (f(x; w) - y)]$$

- And, therefore,

$$\begin{aligned} \nabla^2 e(w) = & E [\nabla^2 f(x; w) (f(x; w) - y)] + \\ & E [\nabla f(x; w) \nabla f(x; w)^T] \end{aligned}$$

- The second term is easy to compute once we have  $\nabla f$
- If  $w \approx w^*$ ,  $f(x; w) \approx y$ ; therefore  $f(x; w) - y \approx 0$ ,
  - We can ignore the first, more complex, term

- We arrive at  $\nabla^2 e(w) \simeq E [\nabla f(x; w) \nabla f(x; w)^\tau]$  or, equivalently,


$$\mathcal{H}_{(i,j)(p,q)}(w) = \left( \frac{\partial^2 e}{\partial w_{ij} \partial w_{pq}}(w) \right) \simeq \left( E \left[ \frac{\partial f}{\partial w_{pq}} \frac{\partial f}{\partial w_{ij}} \right] \right)_{(i,j)(p,q)}$$

- $\mathcal{J} = E [\nabla f(x; w) \nabla f(x; w)^\tau]$  is **Fisher's information matrix**
  - $\mathcal{J}$  is **semidefinite positive**
  - And thus not necessarily invertible
- Often only its diagonal is considered and we have

$$\mathcal{H}_{(i,j)(i,j)}(w) = \left( \frac{\partial f}{\partial w_{ij}} \right)^2$$

- There are many more proposals in unconstrained optimization
- The **Conjugate Gradient** (CG) and **Quasi-Newton** (QN) methods are important in classical MLP training
- The basic idea in CG is to replace gradient descent directions  $g_k = -\nabla e(w_k)$  with new conjugate directions that try to keep somehow the previous “good directions”
- The basic idea in QN is to iterate as in NM but with simple approximations  $\mathcal{A}_k$  to  $\mathcal{H}^{-1}(w^k)$  that converge to  $\mathcal{H}^{-1}(w^*)$
- When training “small” NNs the **Limited-memory Broyden-Fletcher-Goldfarb-Shanno** (L-BFGS) QN variant is often used

- A possibility on the error surface of a NN is to have many long, narrow ravines
  - Gradients bounce in the narrow section, but may be very small along the wider one



- **Momentum** pushes them forward in the wider section (but we have to cope with a new parameter)
- Momentum tries to maintain descent's inertia with a term  $\Delta^k = w^k - w^{k-1}$ , i.e.,

$$w^{k+1} = w^k - \rho_k \nabla_w e(w^k) + \mu_k \Delta^k$$

- More or less, what a **heavy ball** would do
- The goal is to keep  $w^k$  advancing in "plateaux", i.e., small gradient zones
- Momentum can be seen as a crude approximation of a CG step
- Nice explanation at [Why Momentum Really Works](#)

- Let's rewrite momentum in two steps

- 1 Define  $\Delta^{k+1} = -\rho_k \nabla_w e(w^k) + \mu_k \Delta^k$  and
- 2 Apply

$$w^{k+1} = w^k + \Delta^{k+1} = w^k - \rho_k \nabla_w e(w^k) + \mu_k \Delta^k$$

- Nesterov's Accelerated Gradient** is a variant of this

$$\begin{aligned}\tilde{\Delta}^{k+1} &= -\rho_k \nabla_w e(w^k + \mu_k \tilde{\Delta}^k) + \mu_k \tilde{\Delta}^k; \\ w^{k+1} &= w^k + \tilde{\Delta}^{k+1}\end{aligned}$$

- In convex optimization it improves GD and is often highly effective in Deep Network training

- Adam is currently the most widely used gradient-descent method for deep NN training
- At each step  $t$  Adam uses a new random mini-batch to
  - Update **exponentially smoothed** averages  $m_t$  of the gradient  $g_t$  and  $v_t$  of the the squared gradient  $g_t^2 = g_t \odot g_t$  as

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t, \quad v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2;$$

- Since we have

$$E[m_t] \simeq (1 - \beta_1^t) E[g_t], \quad E[v_t] \simeq (1 - \beta_2^t) E[g_t^2],$$

we compute bias corrections  $\hat{m}_t, \hat{v}_t$  as

$$\hat{m}_t = \frac{1}{1 - \beta_1^t} m_t, \quad \hat{v}_t = \frac{1}{1 - \beta_2^t} v_t;$$

- Update weights as  $W_t = W_{t-1} - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$

- Default values  $\alpha = 0.001$ ,  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$ , and  $\epsilon = 10^{-8}$  usually work fine
- What is going on in Adam?
- At each step we work with estimates of the average local gradient:

$$\hat{m}_t \simeq E[\nabla_w f]; \quad \hat{v}_t \simeq \sigma(g_t) \simeq E[(\nabla_w f)^2]$$

- Thus, we can first see the Adam iterations

$$W_t = W_{t-1} - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$$

as a kind of “normalized” gradient descent

- And in more detail, since we can expect  $E[(\nabla_w f)^2] \simeq E[\nabla_{w^2}^2 f]$ , we can see Adam as a “dampened” variant of diagonal Gauss–Newton’s steps



- Typically the  $e(\mathcal{W}^k)$  error diminishes towards an asymptotic minimum
  - If many units are used, we may arrive to 0 error, which usually implies **overfitting**
- First solution: to use a separate **validation subset**  $V$  and stop training when the error in  $V$ , i.e., the **validation error** starts growing
  - This **early stopping** is available in the main NN packages
  - But what do we do for small samples?
- Second solution: **get a good regularization**
  - Now training stops because of reasons such as computational cost, but not because of overfitting risk
- We must combine both, to save on training time and because overfitting is a serious MLP risk

- **Gradient descent** is the simplest method but possibly also the slowest
- **Momentum** can be used to speed it up
- **Newton's method** is much faster but may be very costly and difficult to apply in full form
- All the previous methods require the (usually tricky) selection of a **learning rate**
- Second order methods such as Conjugate Gradient and **Quasi-Newton** avoid learning rates and are more efficient but costlier
- **Limited Broyden-Fletcher-Goldfarb-Shanno** (L-BFGS) is currently the option of choice for "small" MLPs

## 1 Machine Learning Basics

- Basic Modeling
- Regression Basics
- Bias, Variance and Cross Validation
- Basic Classification
- Logistic Regression

## 2 Multilayer Perceptrons

- Rosenblatt's Perceptrons
- Classical MLPs
- Neural Network Classification
- MLP Regularization
- Unconstrained Smooth Optimization
- Computational Costs of MLPs

- Many Vs in Big Data: **Volume**, **Velocity**, Variety, Veracity, Value, ...
- Velocity: information flows in data streams that require fast processing and feed back
  - MLPs are rather fast
  - Less than linear models but much more so than SVMs
- Volume is probably the greatest attractive of Big Data
  - Huge samples and/or very large pattern dimension
  - Large impact in model training

- Parallelism is the first answer to Big Volume
- On a standalone machine it may be
  - Passive: let the SO distribute work among several CPU cores, or use GPUs with low level parallelized libraries such as Linpack or BLAS for linear algebra
  - Active: explicitly exploit a problem's parallelism programming an algorithm in, say, OpenMP
- It is easy to passively parallelize the training of several MLPs
  - But memory costs may multiply
- Training an isolated MLP is not CPU parallelizable:
  - There is a sequential layer dependence in Backprop
  - Each training iteration depends on the previous one
  - But matrix-vector multiplications can exploit GPU parallelism

- Two consequences of a large  $N$  are
  - Sample doesn't fit in memory and we have to split it somehow
  - Training gets "lost", for in the global gradient  $\nabla e = E[\nabla e^\ell]$  we average many local gradients that may cancel each other out
- The first problem has been always present in fields such as analog signal filtering
- Solution for both: adaptive or **on line**, i.e., pattern by pattern, weight updates
  - It also allows a relatively simple setting for a theoretical analysis
  - But not used today; instead, medium-to-large NNs are trained using **mini-batches**

- MLPs are fast to apply but costly to train
- How can we estimate this cost?
- MLP training cost is determined by
  - Its **architecture**, that determines the number of weights to fit and that is also dependent on the input dimension  $D$
  - The full or mini-batch **sample size**  $N$ , that determines the cost of the averages to be computed
  - The **training method**, with more or less iterations that, in turn, are more or less costly

- The number of weights in a single hidden layer (SHL) MLP with  $D$  inputs, one output and one hidden layer with  $H$  units is

$$(D + 1) \times H + (H + 1) \simeq HD$$

- Each extra hidden layer pair with  $(H_1, H_2)$  units adds  $(H_1 + 1) \times H_2 \simeq H_1 H_2$  weights
- The cost per pattern of the feedforward step between layers  $h - 1$ ,  $h$  is basically that of  $H_{h-1} \times H_h$  floating point operations (FPOs)
- Thus, for a general MLP the cost in floating point operations of a forward pass is  $\simeq N \times (\sum_h H_h \times H_{h-1})$ 
  - Very fast on GPUs as it theoretically becomes  $O(N \times \sum_h H_h)$



- Computing a local gradient  $\nabla e^\ell$  in a SHL MLP with square error and one output essentially requires to compute
  - $H$  components for the hidden to output connections, with a  $O(1)$  cost each, for  $\frac{\partial e^\ell}{\partial w_h^O} = (\hat{y} - y)o_h$
  - $DH$  components for the input to hidden connections with essentially an  $O(1)$  cost each, for  $\frac{\partial e^\ell}{\partial w_{h,j}^H} = \left(\sum_{i=1}^L \delta_i w_{ih}^O\right) \sigma'(a_h)x_j$
- In these we are omitting the cost of the partials with respect the bias, but they are much smaller
- Thus, the overall cost of computing  $\nabla e^\ell$  on a SHL MLP is essentially  $O(DH)$
- More layers with  $H_h$  units add a cost  $O(H_{h-1} \times H_h)$ 
  - $O(\sum_h H_h)$  on GPUs

- For a mini-batch of size  $N_b$ , the cost of the mini-batch gradient  $\nabla e$  of a SHL MLP is  $O(N_b \times D \times H)$
- And extra layers add a cost  $O(N_b \times H_{h-1} \times H_h)$
- This is of the same order of magnitude than the cost of the forward pass
- This also dominates the  $O(H_{h-1} \times H_h)$  cost of updating the  $(H_{h-1}, H_h)$  weights in gradient descent
- And all these costs have to be multiplied by the number of training **epochs**

- The important term in the overall training cost is  $\text{nEps} \times \text{cost of } \nabla e$ , with nEps the number of epochs
  - One epoch = one pass on the entire sample
- Thus, the **globally dominant term** in a SHL MLP is

$$\text{nEps} \times \text{cost of } \nabla e = O(\text{nEps} \times N \times D \times H)$$

with  $N$  here the entire sample size

- And for each hidden layer we add an extra cost

$$O(\text{nEps} \times N \times H_{h-1} \times H_h)$$

- Thus, training many layered, large MLPs can be **very costly**
  - But **linear in sample size  $N$**

- MLP complexity is determined by its **architecture**  $\{H_h\}$ , **training procedure** and **sample size**  $N$
- The forward and backward MLP passes have basically the **same complexity**
- Their cost per **pattern and layer** is  $H_{h-1} \times H_h$
- For gradient descent these costs are **multiplied by the number nEps of epochs**
- First order methods essentially **do not add extra complexity**
- Second order methods add extra per iteration costs but should **require less iterations**
- Single MLP training is **not CPU parallelizable**
- But **GPUs** can greatly improve MLP processing costs