

Introduction to Deep Learning

Álvaro Barbero Jiménez
alvaro.barbero.jimenez@gmail.com

Some Deep Learning achievements



The New York Times

Google's AlphaGo Defeats Chinese Go Master in Win for A.I.



HONG KONG — It isn't looking good for humanity.

Transitioning entirely to neural machine translation

Language translation is one of the ways we can give people the power to build community and bring the world closer together. It can help people connect with family members who live overseas, or better understand the perspective of someone who speaks a different language. We use machine translation to translate text in posts and comments automatically, in order to break language barriers and allow people around the world to communicate with each other.




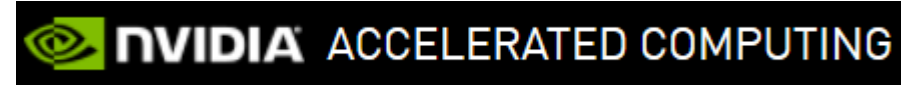
<https://www.nytimes.com/2017/05/23/business/google-deepmind-alphago-go-champion-defeat.html?mcubz=0>

<https://code.facebook.com/posts/289921871474277/transitioning-entirely-to-neural-machine-translation/>

Some Deep Learning achievements


End-to-End Deep Learning for Self-Driving Cars

In a new automotive application, we have used convolutional neural networks (CNNs) to map the raw pixels from a front-facing camera to the steering commands for a self-driving car. This powerful end-to-end approach means that with minimum training data from humans, the system learns to steer, with or without lane markings, on both local roads and highways. The system can also operate in areas with unclear visual guidance such as parking lots or unpaved roads. [Editor's Note: be sure to check out the new post "[Explaining How End-to-End Deep Learning Steers a Self-Driving Car](#)".]




<https://devblogs.nvidia.com/parallelforall/deep-learning-self-driving-cars/>

First generation of neural networks



- Rosenblatt's perceptron (1957)
- Simple learning algorithm
- Solves linear classification problems


First AI Winter



Perceptrons cannot learn
how to solve non-linear
problems


And many practical
problems are non-linear!!

Second generation of neural networks




- Multilayer perceptron (1986)
- Add a layer of hidden units
- More involved learning algorithm (backpropagation)
- Solves non-linear classification problems

Motivation for more layers



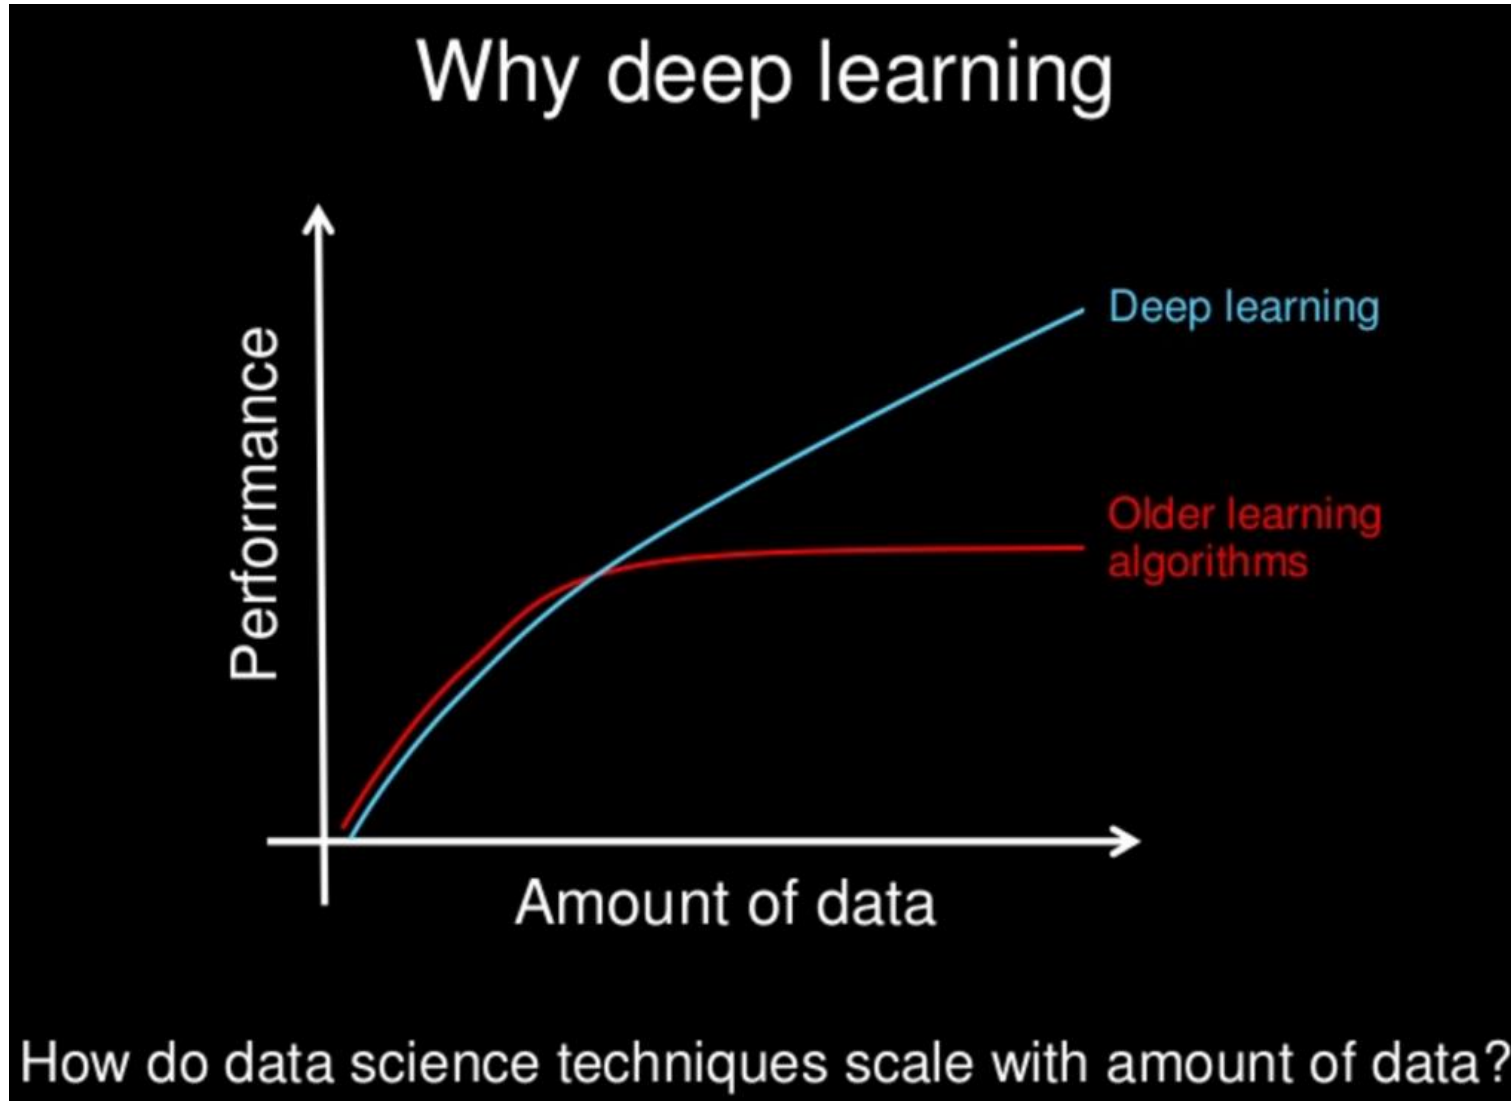
Text-fig. 4. Surface map of a region in post-lateral gyrus. Receptive-field orientations are shown for 32 superficial penetrations. In *A* penetrations are numbered in sequence. In *B* receptive-field orientations are given in degrees of arc, zero corresponding to a horizontal orientation. Interrupted lines are drawn to separate regions of relatively constant axis orientation. Scale, 1 mm. (See also Pl. 2.)

Motivation for more layers




- Visual system as hierarchical structure: several layers of representation
- Each layer builds on top of the previous one to obtain more complex representations
- Mother Nature says: more than one hidden layer can be really useful

Motivation for more layers



Third generation of neural networks




General idea: use **many layers** of hidden units.

→ Deep Learning

In theory, backpropagation can be used to train several hidden layers

Second AI Winter



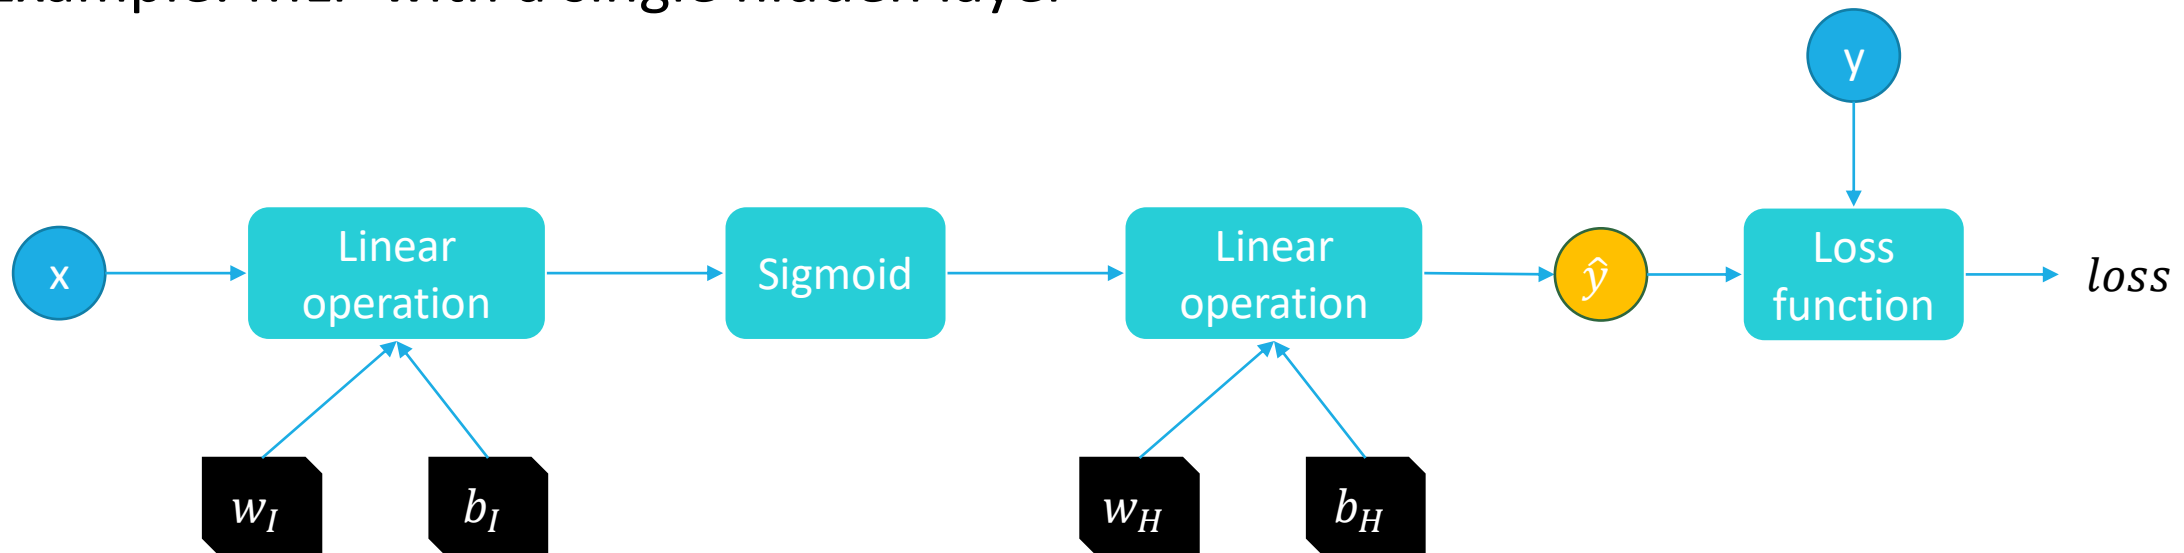
But it doesn't work!

Why??

What is really a neural network?

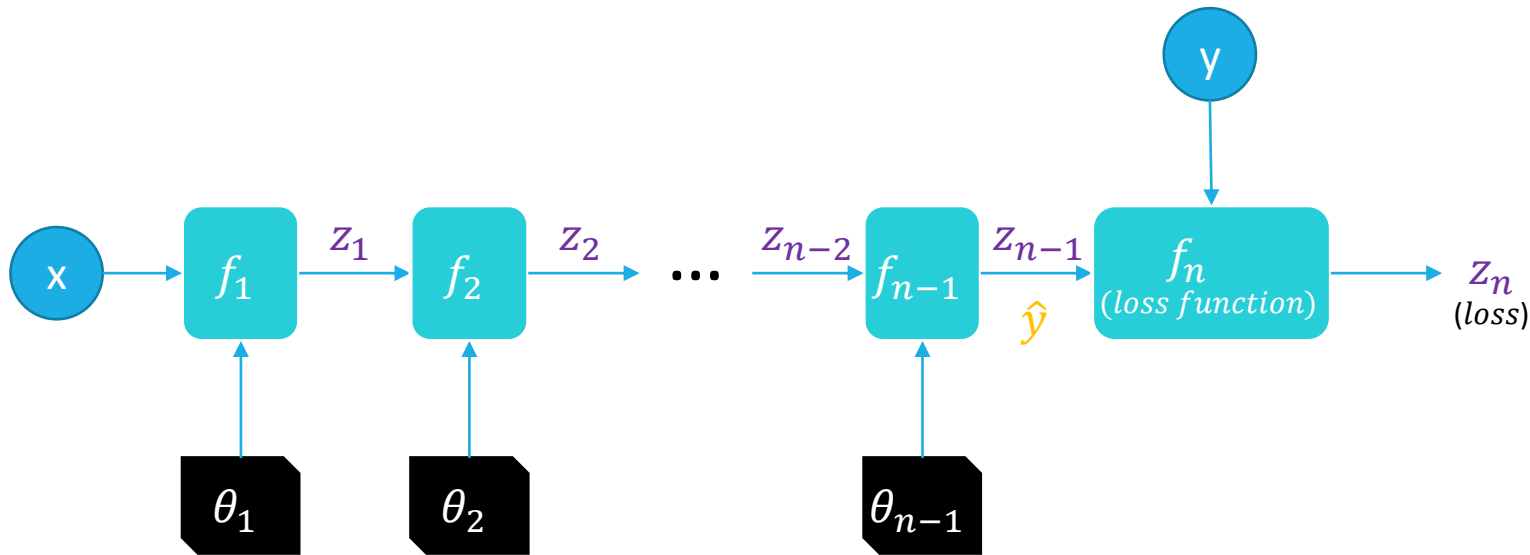
A neural network can be thought as an arbitrary directed acyclic graph (DAG) of **operations**, from an **input vector x** to an **output vector \hat{y}** . Each operation has some parameters to be tuned to minimize a loss function.

Example: MLP with a single hidden layer



What is backpropagation doing?

Take a general network F with layers f_i each with parameters θ_i as




Backpropagation computes the gradients of the whole network for each network parameter. By following the chain rule we have

$$\frac{\partial F}{\partial \theta_i} = \frac{\partial (f_n \circ f_{n-1} \circ \cdots \circ f_1)}{\partial \theta_i} = \frac{\partial z_n}{\partial z_{n-1}} \frac{\partial z_{n-1}}{\partial z_{n-2}} \cdots \frac{\partial z_{i+1}}{\partial z_i} \frac{\partial z_i}{\partial \theta_i} \quad \forall i$$

What is backpropagation doing?

Visually, backpropagation travels the network backwards from $loss$ to θ_i . Each **function** in the path adds a partial derivative of its output with respect to its input.


Example: gradient w.r.t. θ_{n-1} :



$$\frac{\partial F}{\partial \theta_{n-1}} = \frac{\partial z_n}{\partial z_{n-1}} \frac{\partial z_{n-1}}{\partial \theta_{n-1}}$$

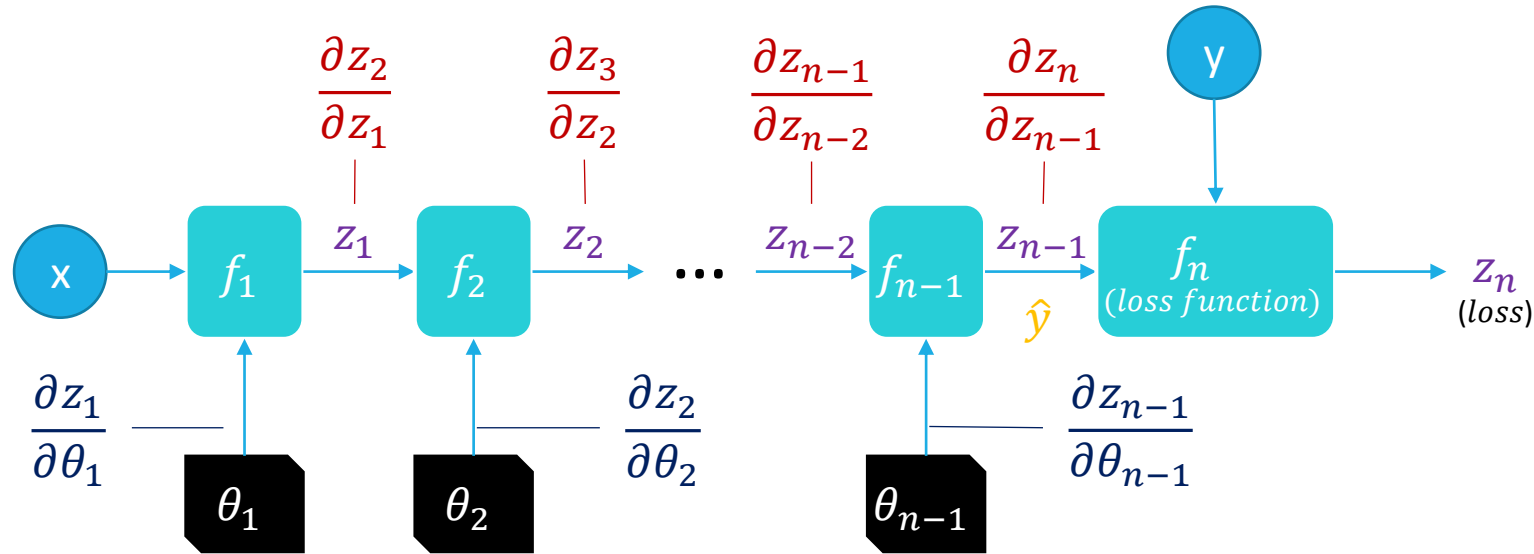
What is backpropagation doing?

Another example: gradient w.r.t. θ_1 :



$$\frac{\partial F}{\partial \theta_1} = \frac{\partial z_n}{\partial z_{n-1}} \frac{\partial z_{n-1}}{\partial z_{n-2}} \dots \frac{\partial z_2}{\partial z_1} \frac{\partial z_1}{\partial \theta_1}$$

Important observations about backpropagation




$$\frac{\partial F}{\partial \theta_{n-1}} = \frac{\partial z_n}{\partial z_{n-1}} \frac{\partial z_{n-1}}{\partial \theta_{n-1}}$$

$$\frac{\partial F}{\partial \theta_{n-2}} = \frac{\partial z_n}{\partial z_{n-1}} \frac{\partial z_{n-1}}{\partial z_{n-2}} \frac{\partial z_{n-2}}{\partial \theta_{n-2}}$$

$$\frac{\partial F}{\partial \theta_{n-3}} = \frac{\partial z_n}{\partial z_{n-1}} \frac{\partial z_{n-1}}{\partial z_{n-2}} \frac{\partial z_{n-2}}{\partial z_{n-3}} \frac{\partial z_{n-3}}{\partial \theta_{n-3}}$$

- Terms in the form $\frac{\partial z_i}{\partial z_{i-1}}$ can be reused throughout $\frac{\partial F}{\partial \theta_i}$ computations $\forall i \rightarrow$ efficient computation of derivatives for all network parameters.
- We can add any function f_i to a neural network, as long as we can compute **every partial derivate** of each output w.r.t. each input.
- Layers closer to the input \rightarrow more terms in gradient calculation

The problem: vanishing gradients

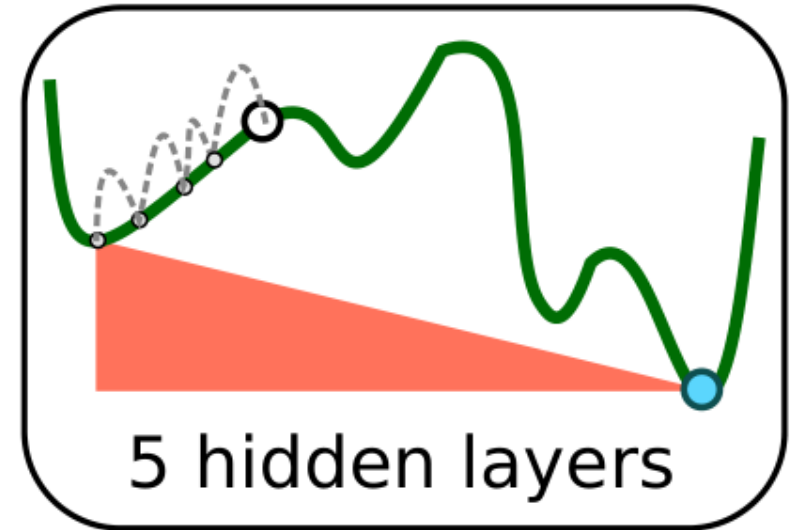
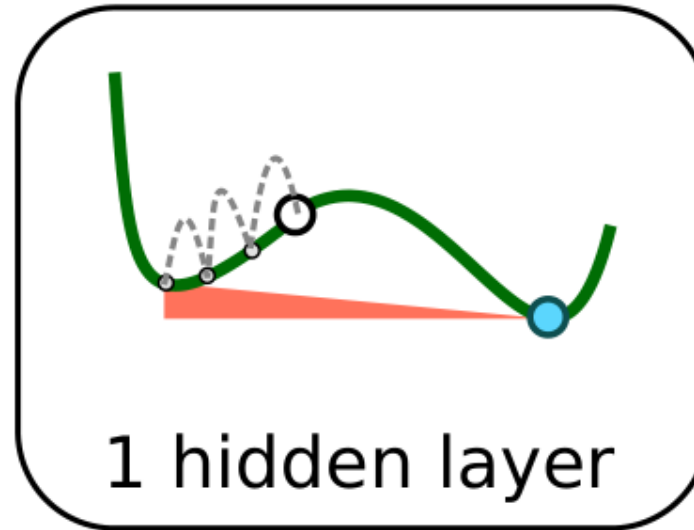



- In backpropagation, the **gradient tends to zero** in the layers nearest to the input
- Intuition: layers near to the output have larger influence in the error rate
- Training deep layers successfully is **terribly slow**
- Until deep layers learn some sensible weights, subsequent layers are trying to **learn mostly from noise**

Glorot & Bengio. Understanding the difficulty in training deep feedforward neural networks

Another problem: high non-linearity

- Backpropagation is a **gradient descent** method.
- Only guarantees finding a **local minimum** of the error function, which can be worse than the **global minimum**
- When introducing more non-linearities in the network (more layers), **local minima multiply**



- It is **hard to find** good parameter configurations for a deep network

How to learn several deep networks

Whole network learning

Use standard neural network methods, but avoid vanishing gradients and other problems by

- Specially designed activations
- Better regularization
- Better backpropagation methods

Present method of choice

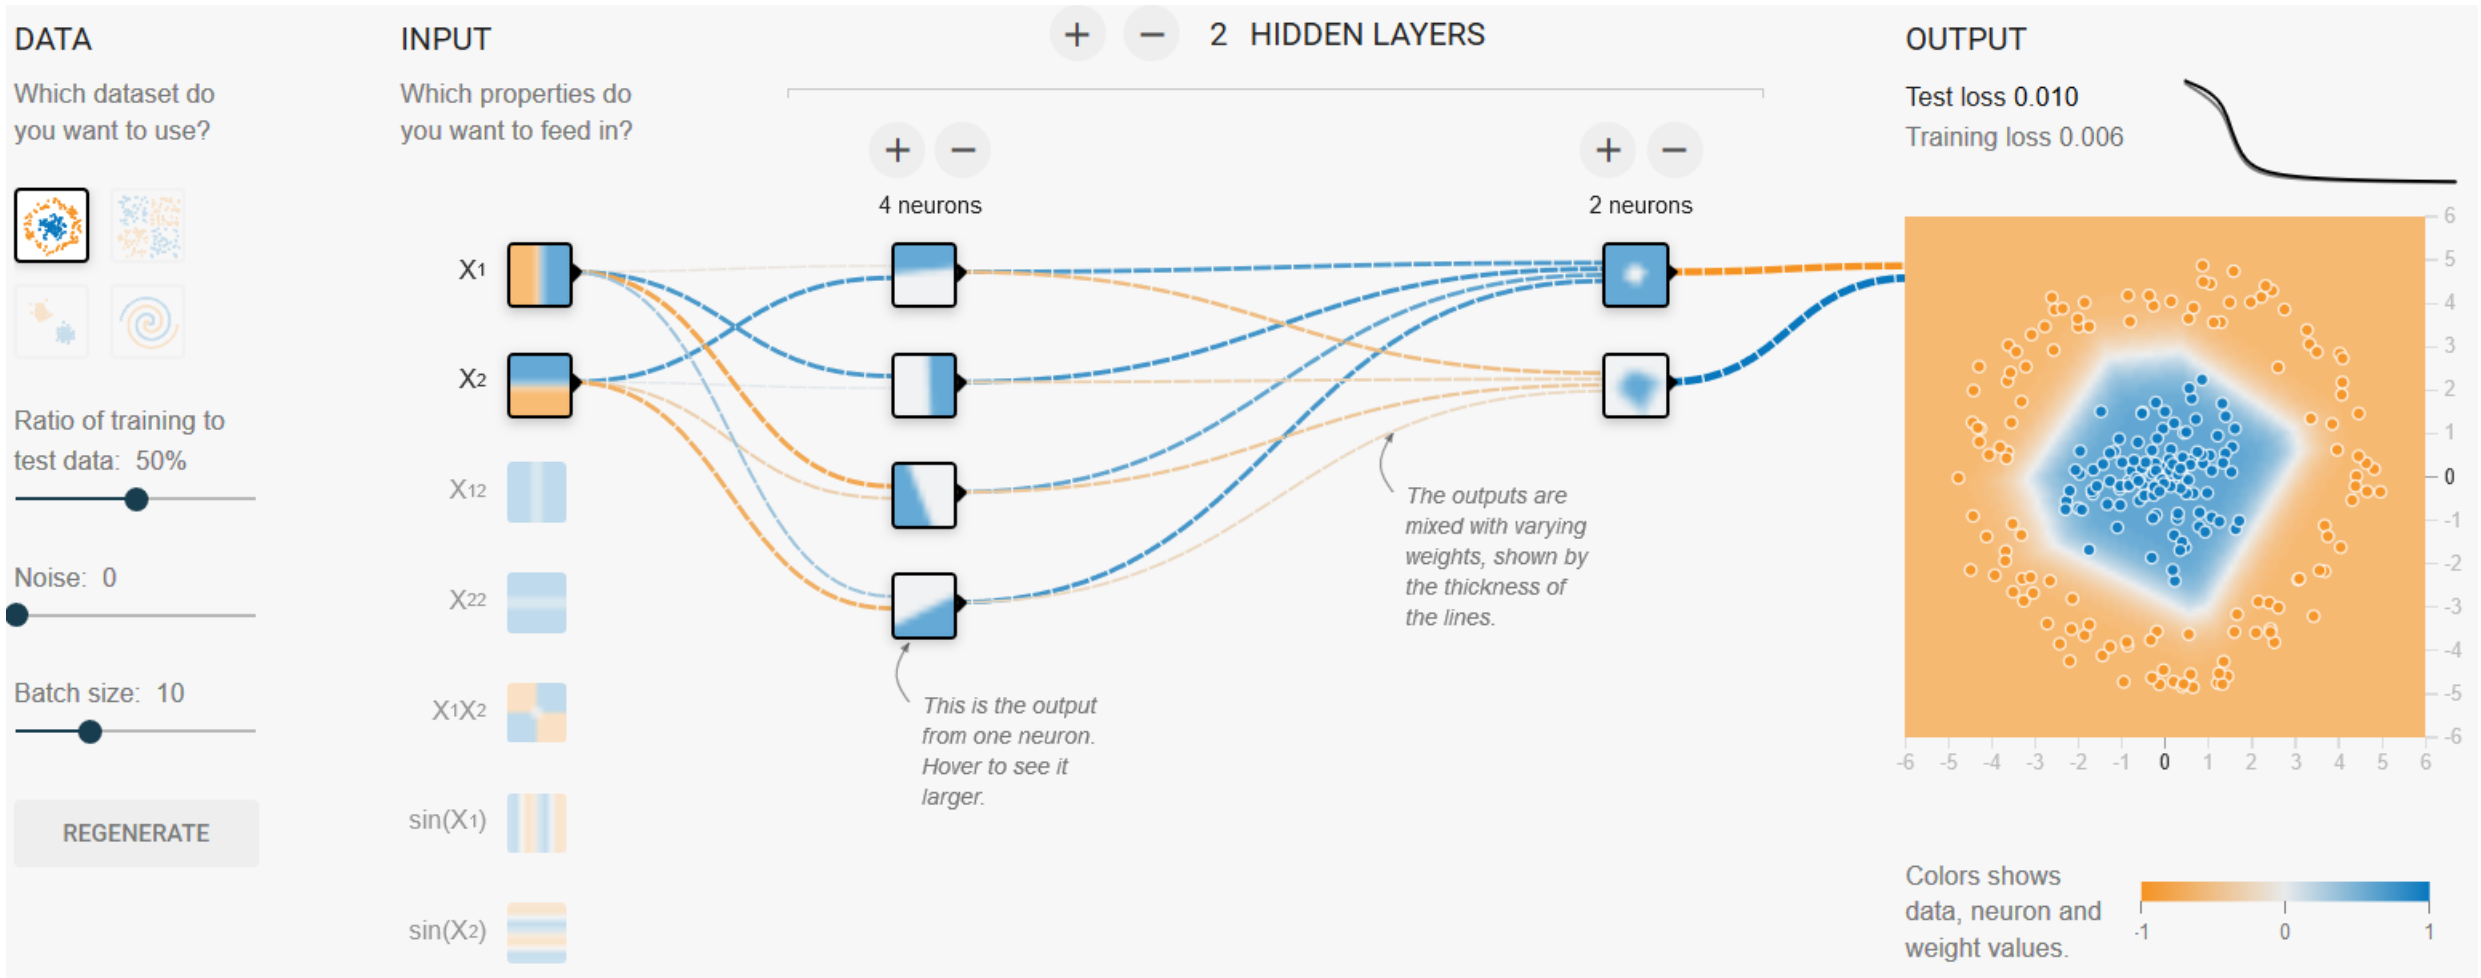
Incremental learning

Learn layers one by one, in an iterative fashion

- Deterministic approach:
Autoencoders
- Probabilistic approach:
Restricted Boltzmann Machines

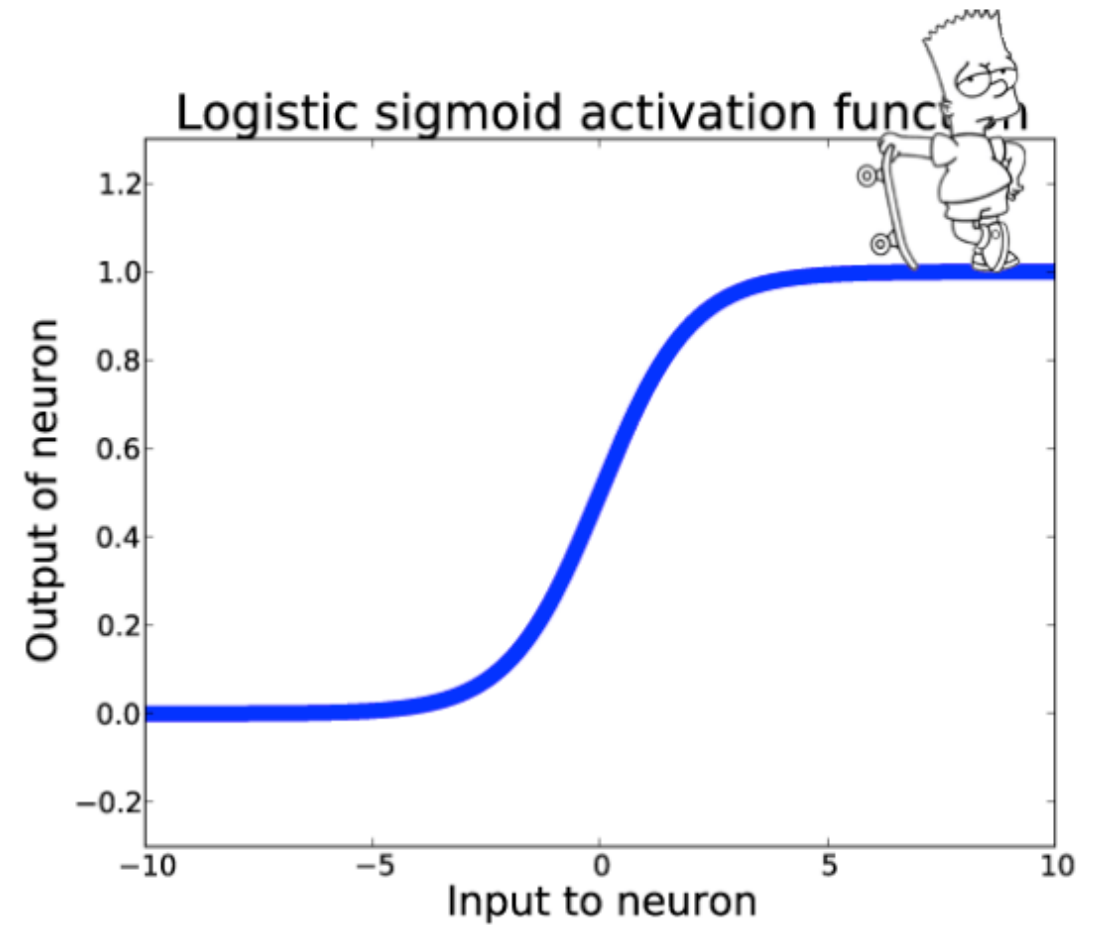

Popular in the first years of Deep Learning, now less used

Working example



A Neural Network playground - <http://playground.tensorflow.org>

The Bart Simpson problem




Goodfellow - Deep learning of representations and its application to computer vision

[Serengil – An overview of the Vanishing Gradient Problem](#)

The Simpsons by Matt Groening

The Bart Simpson problem in backpropagation terms




If the network includes a layer that can produce ≈ 0 terms in the partial derivative, all parameters in previous layers are likely to get very small gradients.

$$\frac{\partial F}{\partial \theta_i} = \frac{\partial z_n}{\partial z_{n-1}} \frac{\partial z_{n-1}}{\partial z_{n-2}} \dots \frac{\partial z_{i+1}}{\partial z_i} \frac{\partial z_i}{\partial \theta_i} = \frac{\partial z_n}{\partial z_{n-1}} (\approx 0) \dots \frac{\partial z_{i+1}}{\partial z_i} \frac{\partial z_i}{\partial \theta_i} \approx 0 \quad \forall i$$

Rectified Linear Units

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$



- A better choice for the activation function is to use **Rectified Linear Units**
- Very cheap to compute (no products or exponentials)
- Faster training
- Leads to sparser neuron activations

!!! Non-differentiable

- Backpropagation is performed by computing the derivative of f as



$$\frac{\partial f}{\partial x} \simeq \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$

- This is a form of **subgradient descent**
- **No small gradients in this formula**

Zeiler et al. On Rectified Linear Units For Speech Processing.

Leaky and Parametric Rectified Linear Units

Some papers show that including a small slope in the negative side of ReLU units (Leaky ReLU) can improve model performance. This slope parameter a can also be learned by the network (PReLU).



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ ax & \text{if } x \leq 0 \end{cases}$$

$$\frac{\partial f}{\partial x} \simeq \begin{cases} 1 & \text{if } x > 0 \\ a & \text{if } x \leq 0 \end{cases}$$

Maas et al. Rectified Nonlinearities Improve Neural Network Acoustic Models

He et al. Delving Deeper into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification

Xu et al. Empirical Evaluation of Rectified Activations in Convolution Network

[Serengil – An overview of the Vanishing Gradient Problem](#)

The Simpsons by Matt Groening

Activation	Training Error	Test Error
ReLU	0.00318	0.1245
Leaky ReLU, $a = 100$	0.0031	0.1266
Leaky ReLU, $a = 5.5$	0.00362	0.1120
PReLU	0.00178	0.1179
RReLU ($y_{ji} = x_{ji}/\frac{l+u}{2}$)	0.00550	0.1119

Table 3. Error rate of CIFAR-10 Network in Network with different activation function


Activations and losses in output layer

Problem class Output layer activation


Loss function

Binary or
multilabel
classification

Sigmoid $f(x_i) = \frac{1}{1+e^{-x_i}}$



Binary cross entropy $l(\hat{y}, y) = -(y \log(\hat{y}) + (1 - y) \log(1 - \hat{y}))$



Multiclass
classification

Softmax $f(x_i) = \frac{e^{x_i}}{\sum_d e^{x_d}}$

Categorical cross entropy

$$l(\hat{y}, y) = - \sum_c (y_c \log(\hat{y}_c))$$

With y one-hot vector encoding true label, \hat{y} output vector of class probabilities

Regression

Linear $f(x_i) = x_i$

Absolute error $l(\hat{y}, y) = |\hat{y} - y|$
Squared error $l(\hat{y}, y) = \|\hat{y} - y\|_2^2$

Regularization through Dropout

- For each pattern introduced in the network, **randomly deactivate units with probability p**
- Implicitly **trains 2^n networks sharing the same weights**, but with all the possible combinations of the activated/deactivated n neurons of the network
- Forces the network develop neurons robust to missing connections





Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.


Srivastava. Improving Neural Networks with Dropout

Srivastava et al. Dropout: A simple way to prevent neural networks from overfitting


Dropout as a deep network layer



Regularization through Dropout - Prediction



Training time



Prediction time

- At prediction time neurons are never dropped, but to compensate we need to scale each weight in the network by the inverse dropout probability $1-p$
- With this scaling we approximate the average of the 2^n implicitly trained networks

Srivasta. Improving Neural Networks with Dropout


Srivastava et al. Dropout: A simple way to prevent neural networks from overfitting

Regularization through Dropout - Performance

Method	Unit Type	Architecture	Error %
Standard Neural Net (Simard et al., 2003)	Logistic	2 layers, 800 units	1.60
SVM Gaussian kernel	NA	NA	1.40
Dropout NN	Logistic	3 layers, 1024 units	1.35
Dropout NN	ReLU	3 layers, 1024 units	1.25
Dropout NN + max-norm constraint	ReLU	3 layers, 1024 units	1.06
Dropout NN + max-norm constraint	ReLU	3 layers, 2048 units	1.04
Dropout NN + max-norm constraint	ReLU	2 layers, 4096 units	1.01
Dropout NN + max-norm constraint	ReLU	2 layers, 8192 units	0.95
Dropout NN + max-norm constraint (Goodfellow et al., 2013)	Maxout	2 layers, (5 × 240) units	0.94
DBN + finetuning (Hinton and Salakhutdinov, 2006)	Logistic	500-500-2000	1.18
DBM + finetuning (Salakhutdinov and Hinton, 2009)	Logistic	500-500-2000	0.96
DBN + dropout finetuning	Logistic	500-500-2000	0.92
DBM + dropout finetuning	Logistic	500-500-2000	0.79

Table 2: Comparison of different models on MNIST.

The MNIST data set consists of 28×28 pixel handwritten digit images. The task is to classify the images into 10 digit classes. Table 2 compares the performance of dropout with other techniques. The best performing neural networks for the permutation invariant



Srivastava et al. Dropout: A simple way to prevent neural networks from overfitting

Regularization through Dropout – Learned units

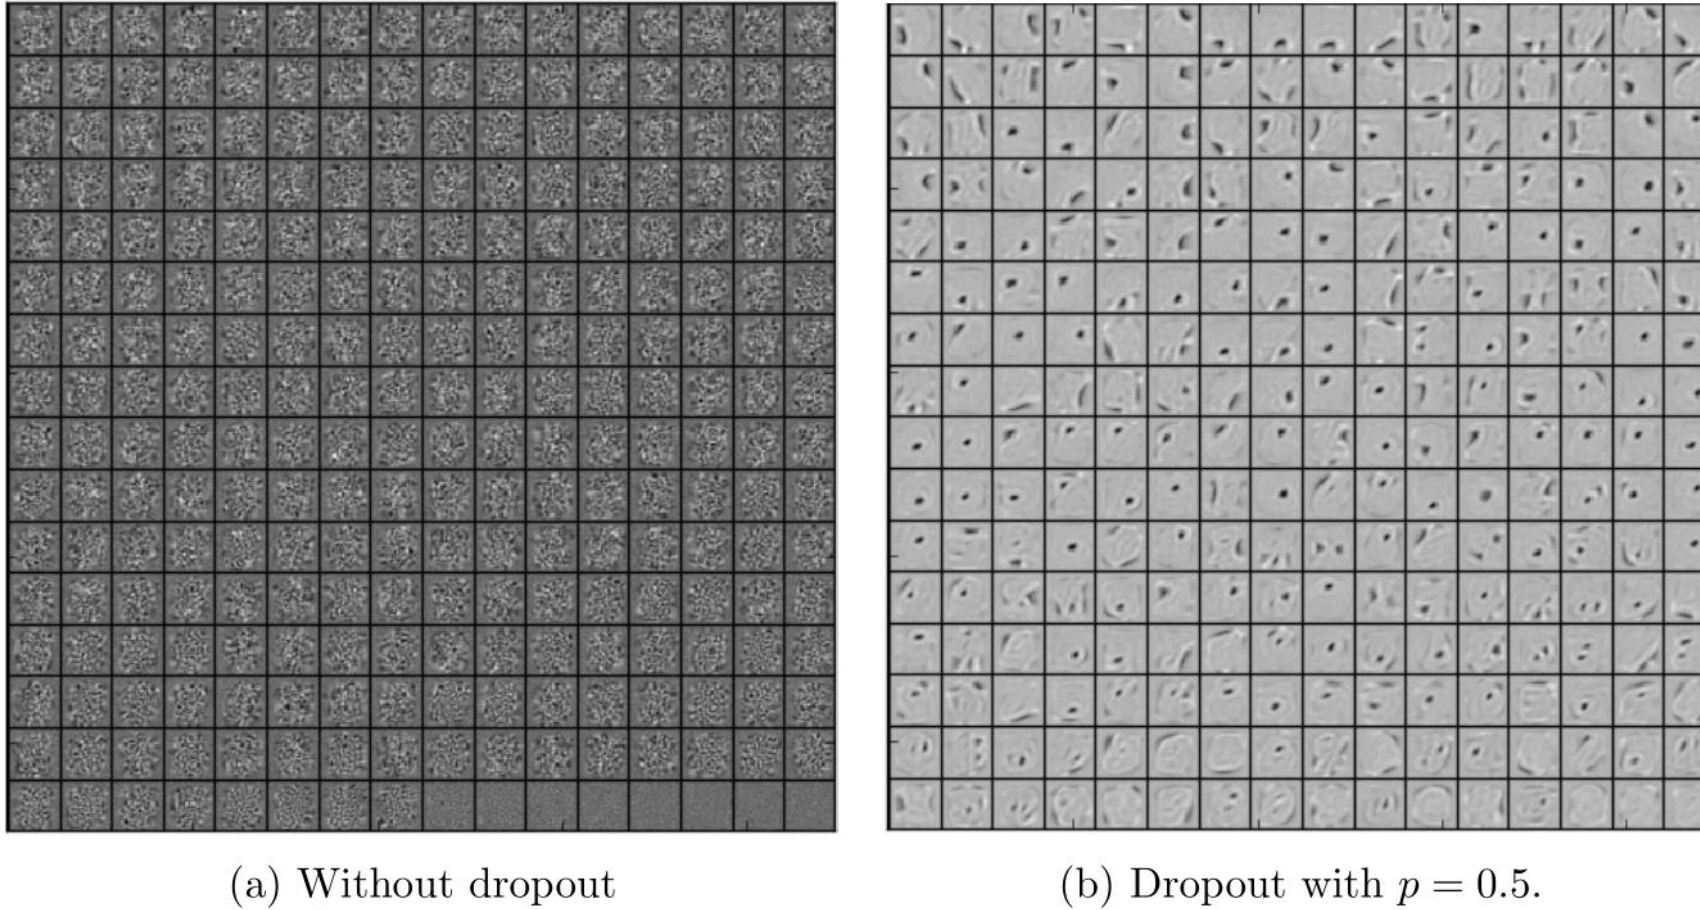



Figure 7: Features learned on MNIST with one hidden layer autoencoders having 256 rectified linear units.

Designing Deep learning models




- Think in layers, not neurons
 - Each layer is stacked on top of others to achieve higher levels of abstraction
- Different kinds of layers are used
 - Data input layers
 - Output (loss) layers
 - Fully connected (“classic”) layers
 - Convolutional layers for image processing
 - Recurrent layers for sequence processing
 - ...

Stacking layers

Trend in the last years:
leverage advances in
optimization methods,
faster hardware and
network design to stack
more and more layers
of neurons.

ImageNet competition:
deeper networks every
year!



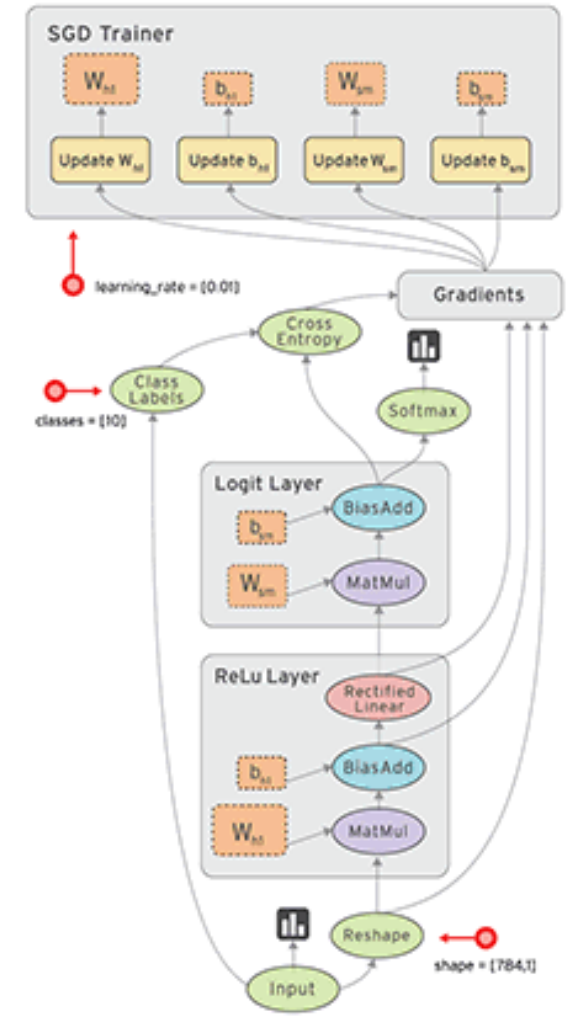
IMAGENET



- 2012: AlexNet, 8 layers ([Krizhevsky et al. – ImageNet Classification with Deep Convolutional Neural Networks](#))
- 2014: VGG, 19 layers ([Simonyan & Zisserman – Very Deep Convolutional Networks for Large-Scale Image Recognition](#))
- 2015: ResNet - 152 layers ([He et al. – Deep Residual Learning for Image Recognition](#))
- 2017: Squeeze-and-excitation – 152 layers ([Hu et al. – Squeeze-and-Excitation Networks](#))

Modern approach to building deep networks

1. Define network structure: computational graph (structure), layer types (functions), loss
2. Use a automatic differentiation backend to infer backpropagation calculations
 - The backend automatically generates C/CUDA code
3. Train the network with an optimization method



Deep Learning frameworks

High-level

(network design)



keras

PYTORCH



Low-level

(symbolic differentiation
backend)



theano
(deprecated)



dmlc
mxnet

Example: learning in TensorFlow



TensorFlow - <https://www.tensorflow.org/>

```
import tensorflow as tf
import numpy as np

# Create 100 phony x, y data points in NumPy, y = x * 0.1 + 0.3
x_data = np.random.rand(100).astype(np.float32)
y_data = x_data * 0.1 + 0.3

# Try to find values for W and b that compute y_data = W * x_data + b
# (We know that W should be 0.1 and b 0.3, but Tensorflow will
# figure that out for us.)
W = tf.Variable(tf.random_uniform([1], -1.0, 1.0))
b = tf.Variable(tf.zeros([1]))
y = W * x_data + b

# Minimize the mean squared errors.
loss = tf.reduce_mean(tf.square(y - y_data))
optimizer = tf.train.GradientDescentOptimizer(0.5)
train = optimizer.minimize(loss)

# Before starting, initialize the variables. We will 'run' this first.
init = tf.initialize_all_variables()

# Launch the graph.
sess = tf.Session()
sess.run(init)

# Fit the line.
for step in xrange(201):
    sess.run(train)
    if step % 20 == 0:
        print(step, sess.run(W), sess.run(b))

# Learns best fit is W: [0.1], b: [0.3]
```

Example: learning in Keras




```
# For a single-input model with 2 classes (binary classification):

model = Sequential()
model.add(Dense(32, activation='relu', input_dim=100))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['accuracy'])


# Generate dummy data
import numpy as np
data = np.random.random((1000, 100))
labels = np.random.randint(2, size=(1000, 1))

# Train the model, iterating on the data in batches of 32 samples
model.fit(data, labels, epochs=10, batch_size=32)
```



Physical reading material



<https://www.deeplearningbook.org/>



<https://www.amazon.es/Deep-Learning-Python-Francois-Chollet/dp/1617294438>



<https://www.amazon.com/Neural-Networks-Babies-Baby-University/dp/1492671207>



Afi Escuela

© 2021 Afi Escuela. Todos los derechos reservados.